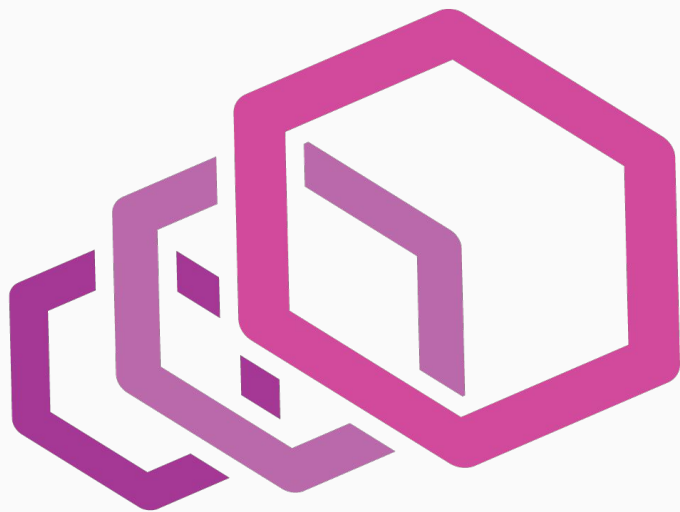


envoy

Observability and control in the age of the service mesh: present and future

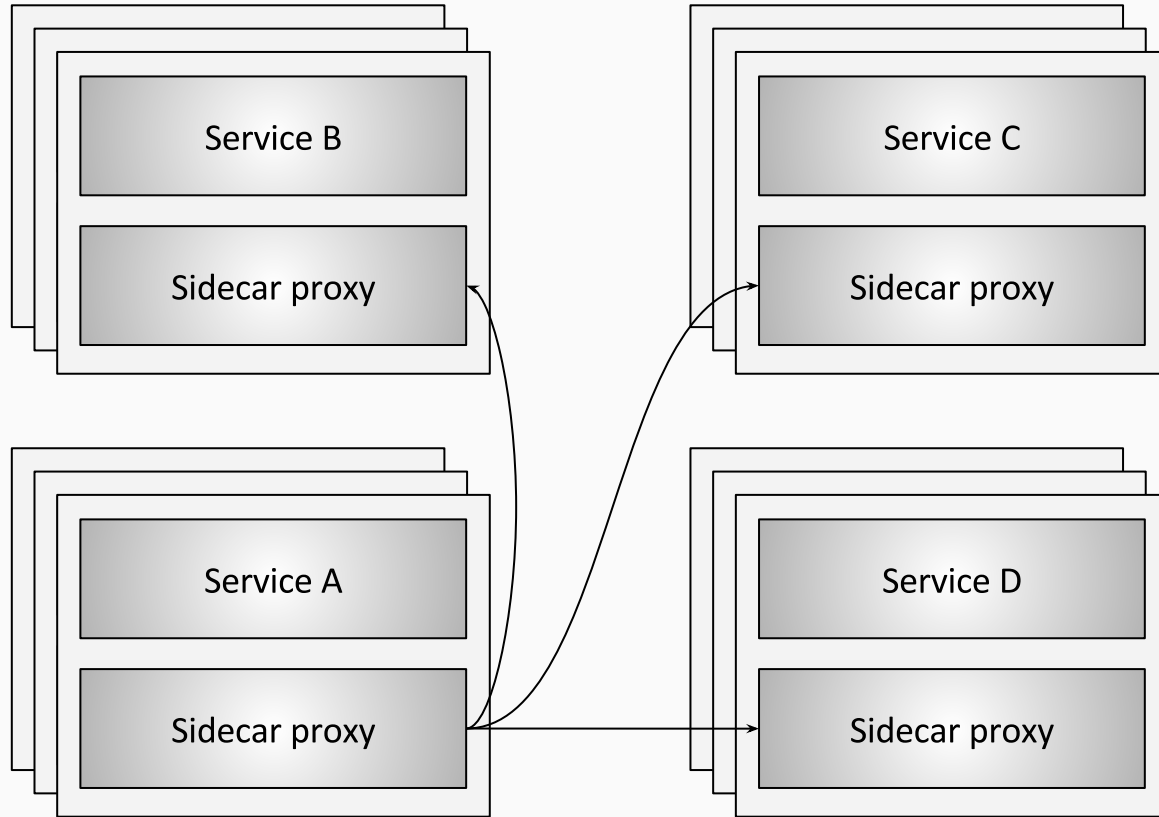
What is Envoy and the service mesh?

The network should be transparent to applications. When network and application problems do occur it should be easy to determine the source of the problem.



envoy

Service mesh refresher



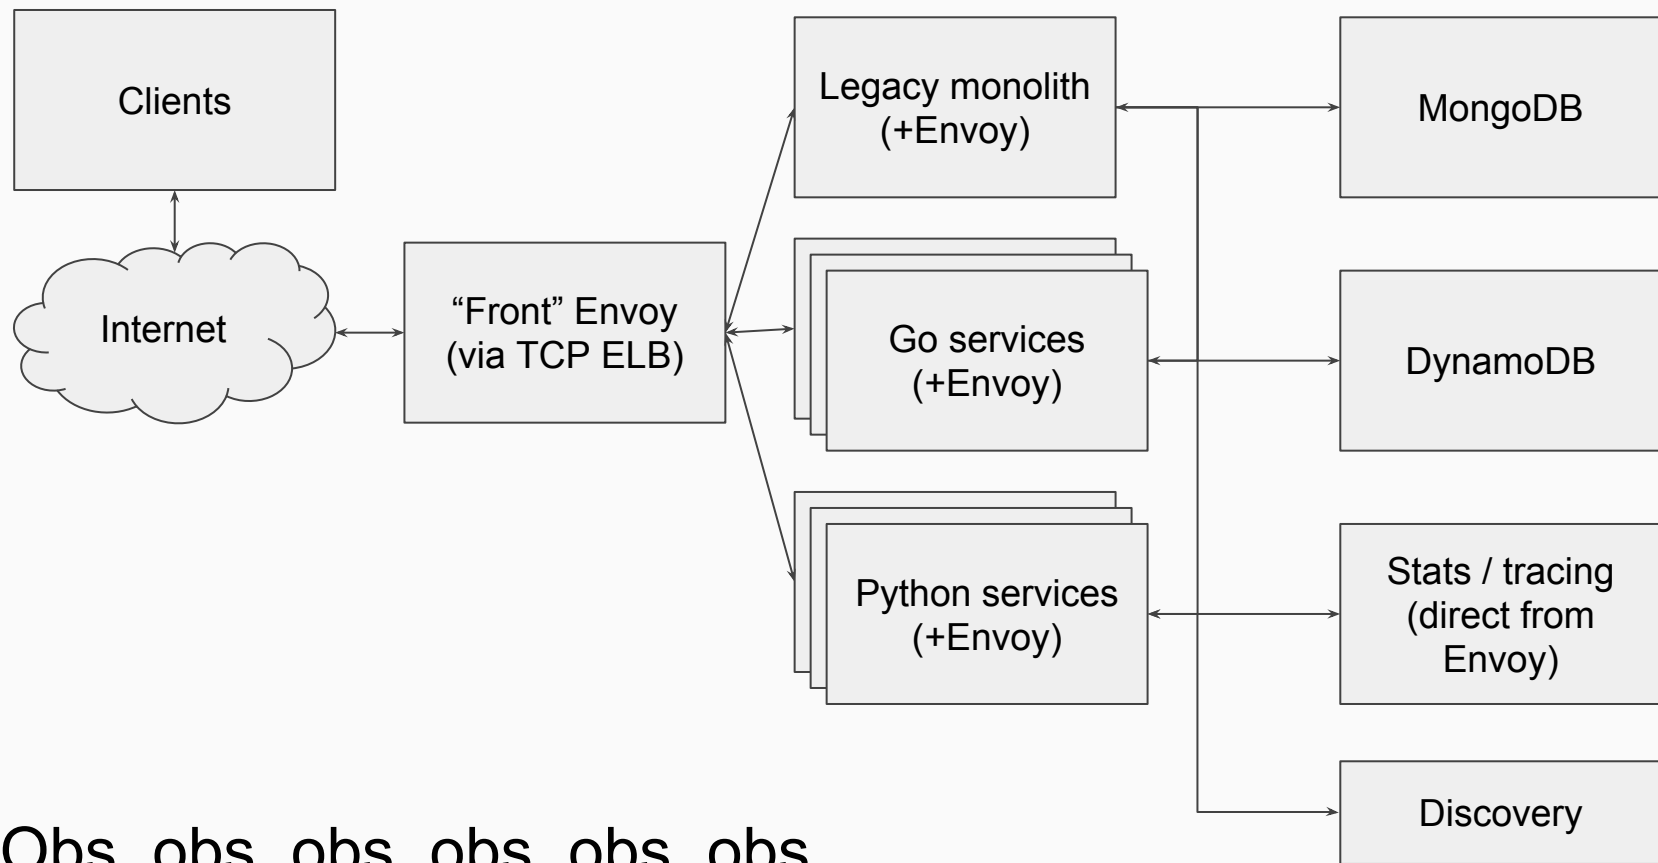
- **Out of process architecture**
- **Modern C++11 code base**
- **L3/L4 filter architecture**
- **HTTP L7 filter architecture**
- **HTTP/2 first**
- **Service discovery and active/passive health checking**
- **Advanced load balancing**
- **Best in class observability** (stats, logging, and tracing)
- **Edge proxy**

Observability

- **Observability** is by far the most important thing that Envoy provides.
- Having all SoA traffic transit through Envoy gives us a single place where we can:
 - Produce consistent **statistics** for every hop
 - Create and propagate a stable **request ID / tracing context**
 - Consistent **logging**
 - Distributed **tracing**



Lyft today



Obs, obs, obs, obs, obs, obs...

The page goes out (hopefully). What is the best case scenario of what follows?



State of incident handling @lyft: the page

[Incidents](#)[Alerts](#)[Configuration ▾](#)[Analytics ▾](#)[Add-ons ▾](#)[Command Console](#)

Your open incidents

0 triggered 0 acknowledged

All open incidents

51 triggered 42 acknowledged

Open **Triggered** Acknowledged Resolved Any Status

! Acknowledge

🔄 Reassign

✓ Resolve

⌚ Snooze ▾

<input type="checkbox"/>	Status	Urgency ▾	Title ▾	Created ▾	Service
<input type="checkbox"/>	Triggered	High	[STAGING] developerportal / Autoscaling / Envoy Membership / Percent Healthy Hosts (developerportal) <small>🔒 HIDE DETAILS #131276</small>	on Jul 31, 2017 at 9:12 AM	developerportal-low-urgency

CUSTOM DETAILS

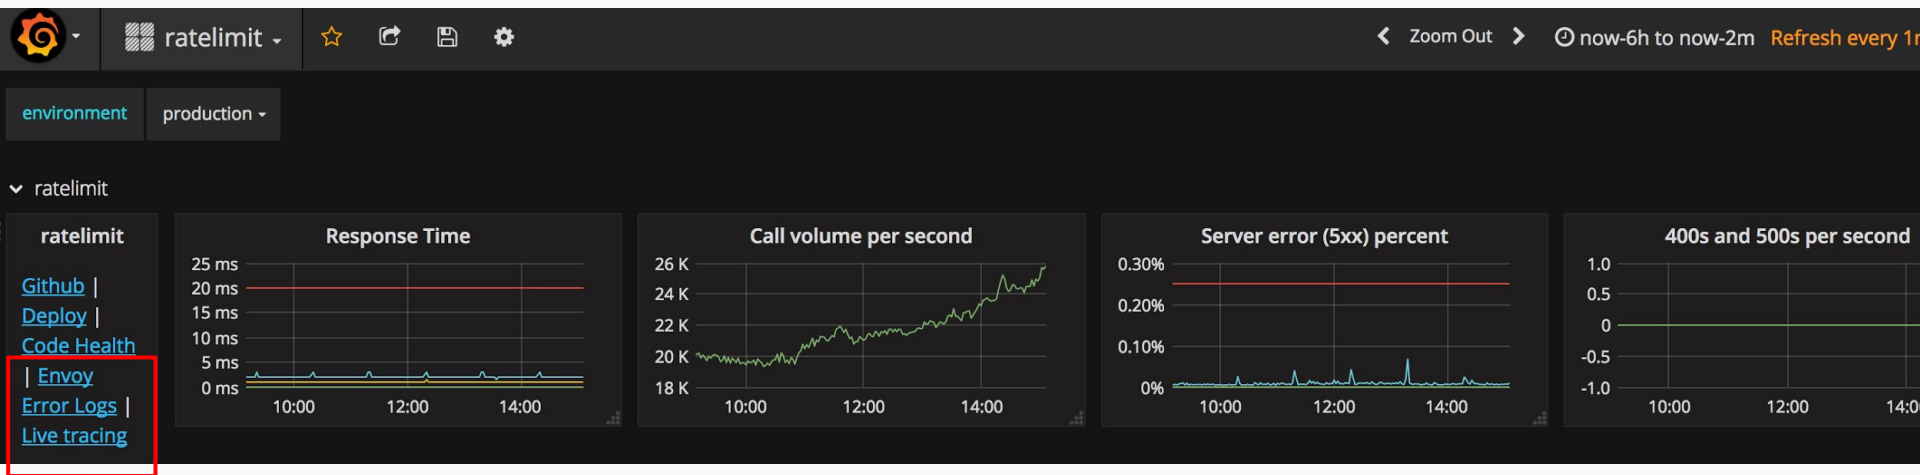
Environment	staging
Dashboard	https://grafana.lyft.net/dashboard/db/developerportal?var-environment=staging
Condition	<code>(rawavg(aligned(1m, percentile(20, ts(staging.infra.aws.ec2.asg.envoy.cluster.developerportal.membership_healthy.gauge.mean)))) / rawavg(aligned(1m, percentile(20, ts(staging.infra.aws.ec2.asg.envoy.cluster.developerportal.membership_total.gauge.mean)))) * 100) < 85</code>
Alert	https://lyft.wavefront.com/alerts/1500501889631/

CLIENT

[View in wavefront](#)

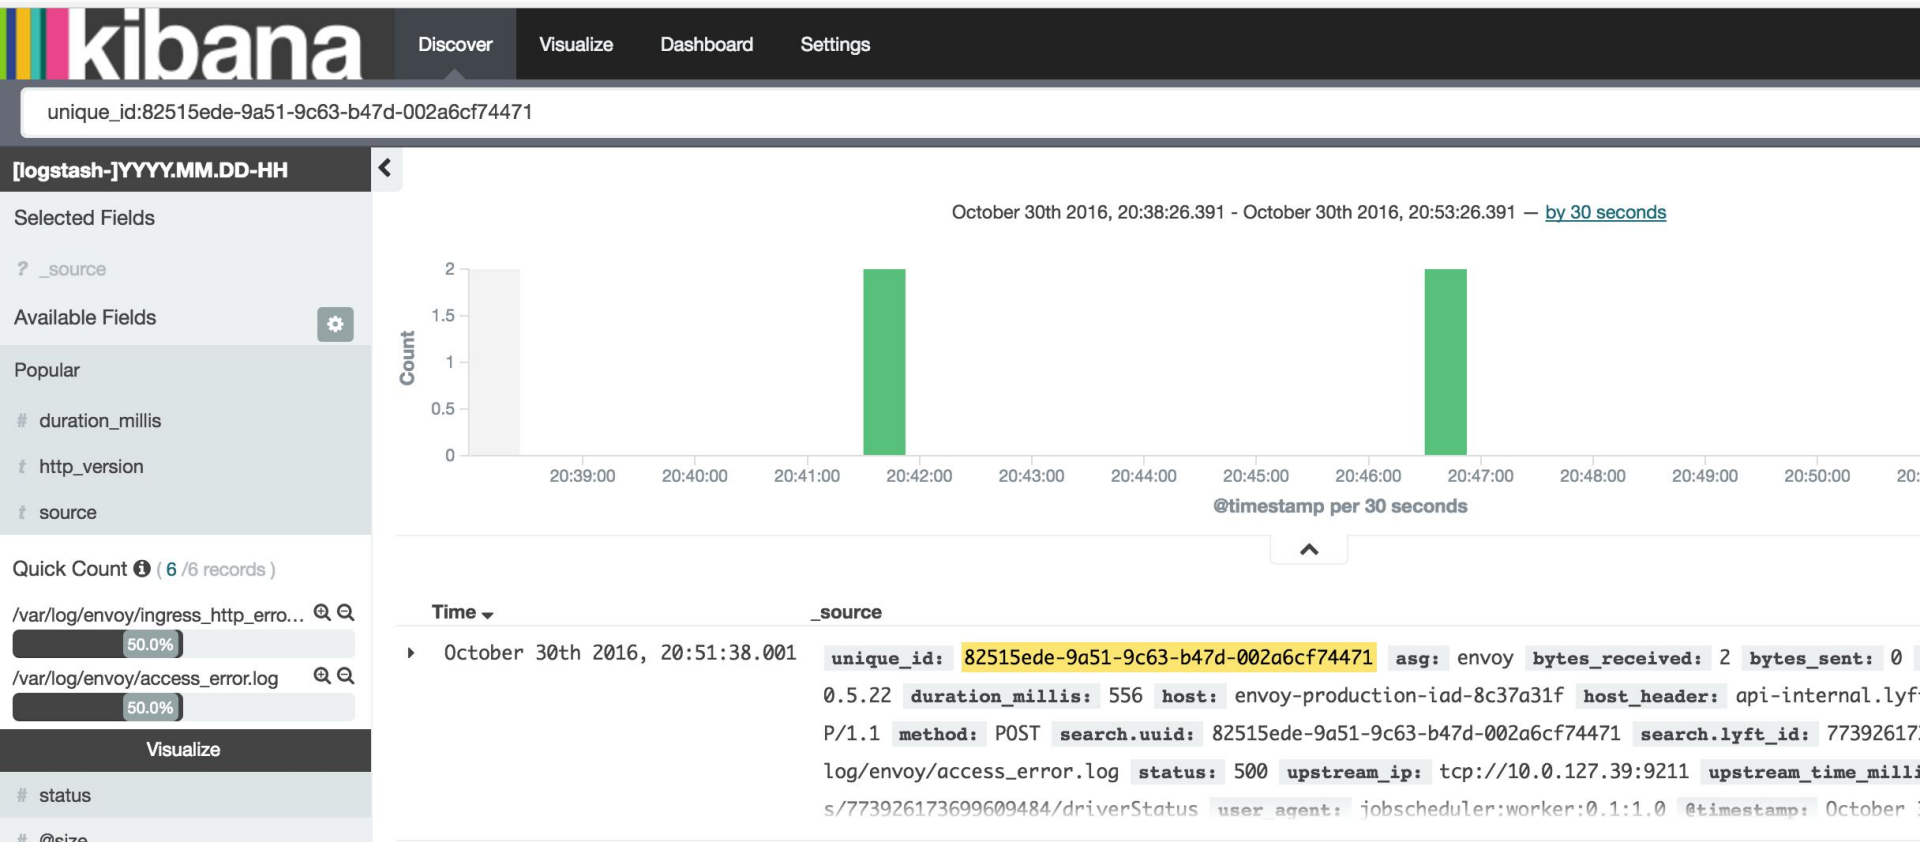
[View Message](#)

State of incident handling @lyft: per service auto-generated panel



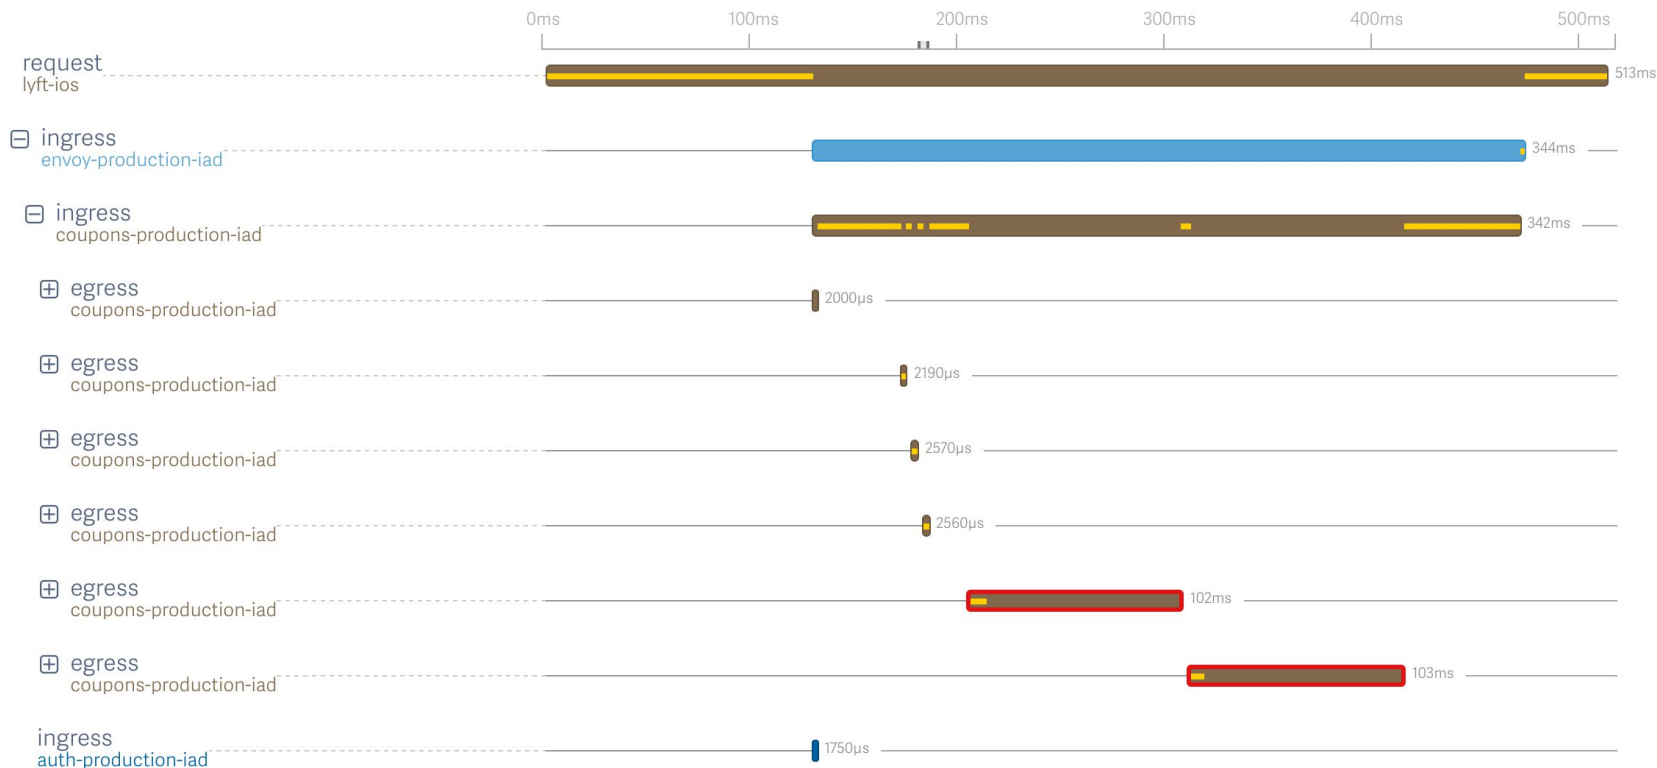
Links to logging and tracing

State of incident handling @lyft: logging



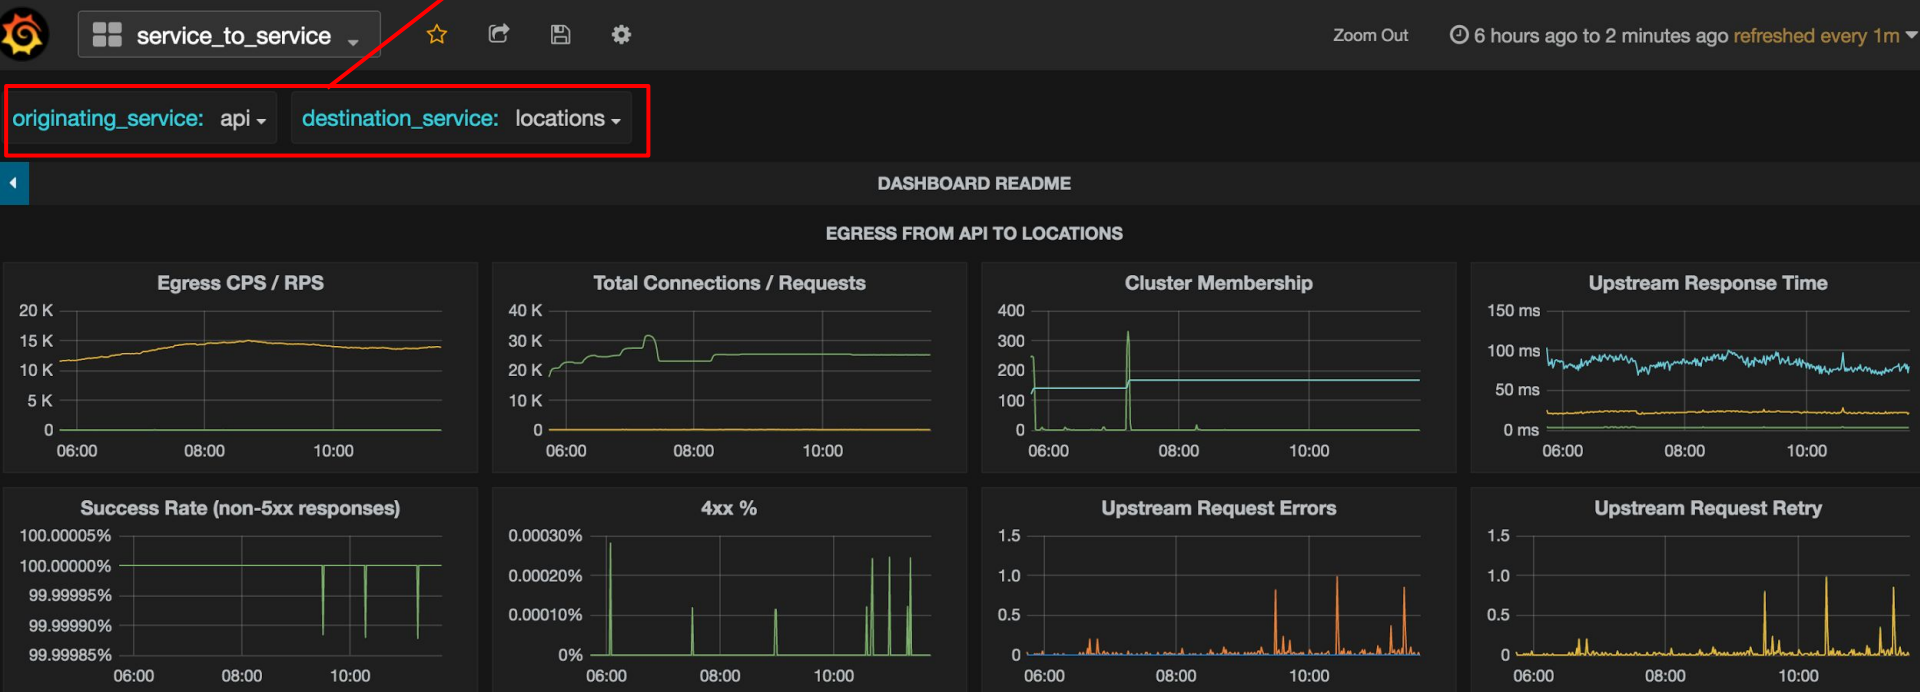
State of incident handling @lyft: distributed tracing

Expand all spans



State of incident handling @lyft: service to service template dashboard

Template with drop down for every service



State of incident handling @lyft: edge proxy

▼ PER HOST

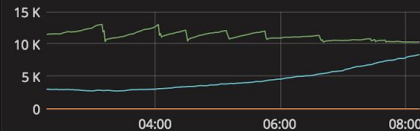
Downstream CPS



Downstream RPS



Downstream Total Connections



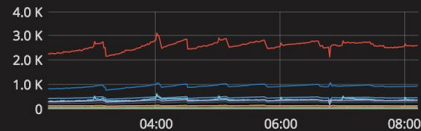
Downstream Total Requests



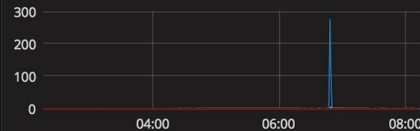
Downstream Connection Length



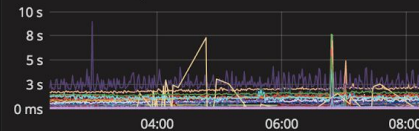
Upstream RPS All Clusters



Upstream 5xx All Clusters



Upstream P99 All Clusters



State of incident handling @lyft: global health dashboard

 Envoy-Global

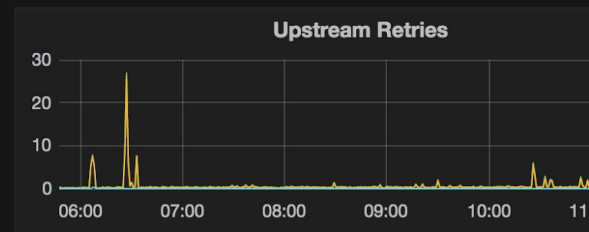
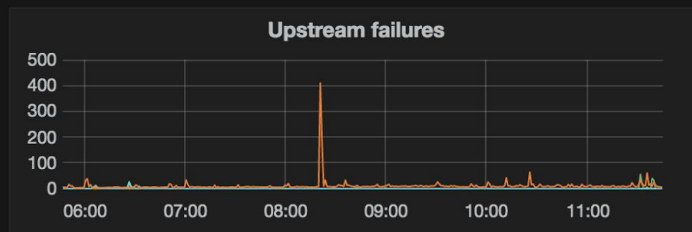
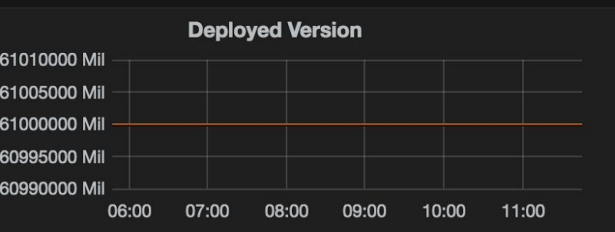
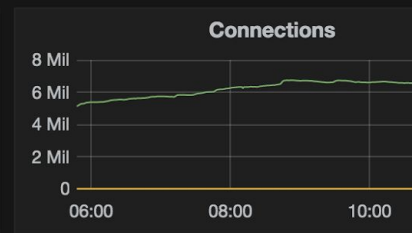
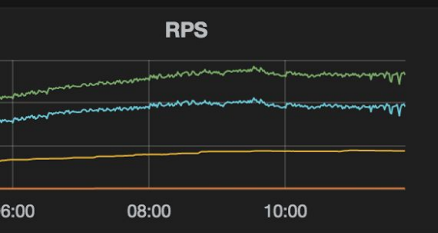


Zoom Out

🕒 6 hours ago to 2 minutes ago refreshed ev

DASHBOARD README

TOP LEVEL ALL ENVOYS



CROSS ZONE TRAFFIC

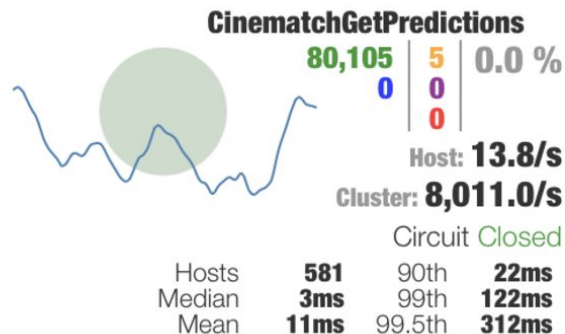
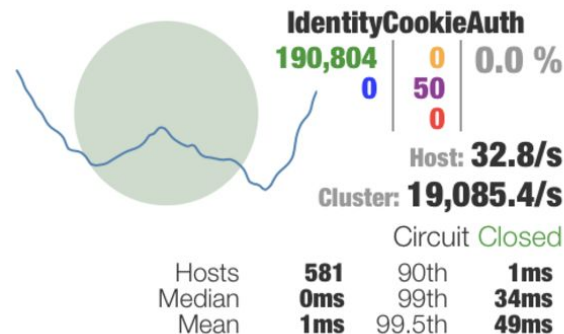
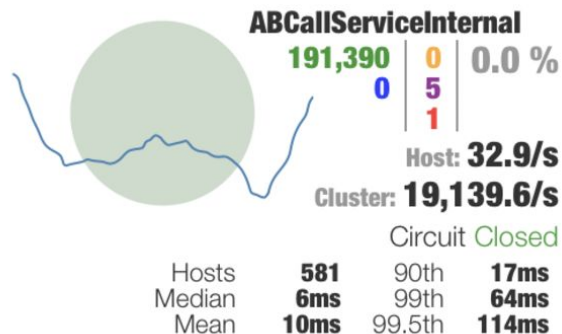
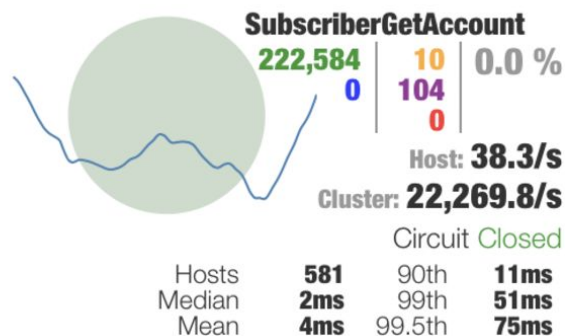
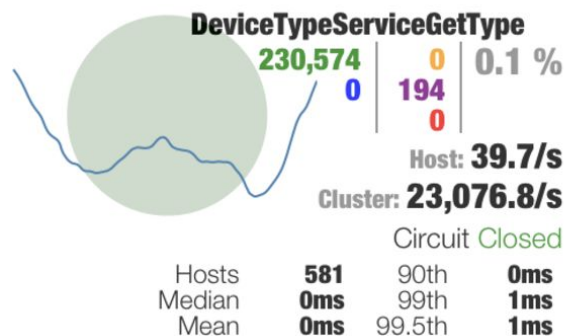
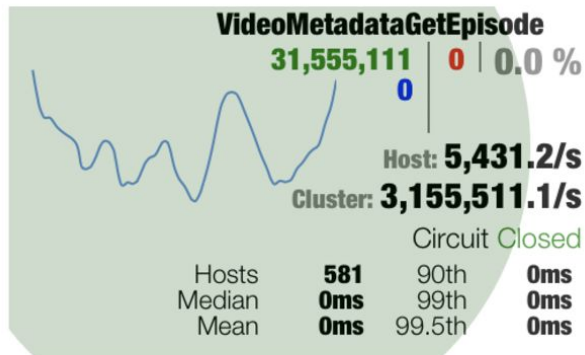
RATELIMIT

Future of microservice observability: problems

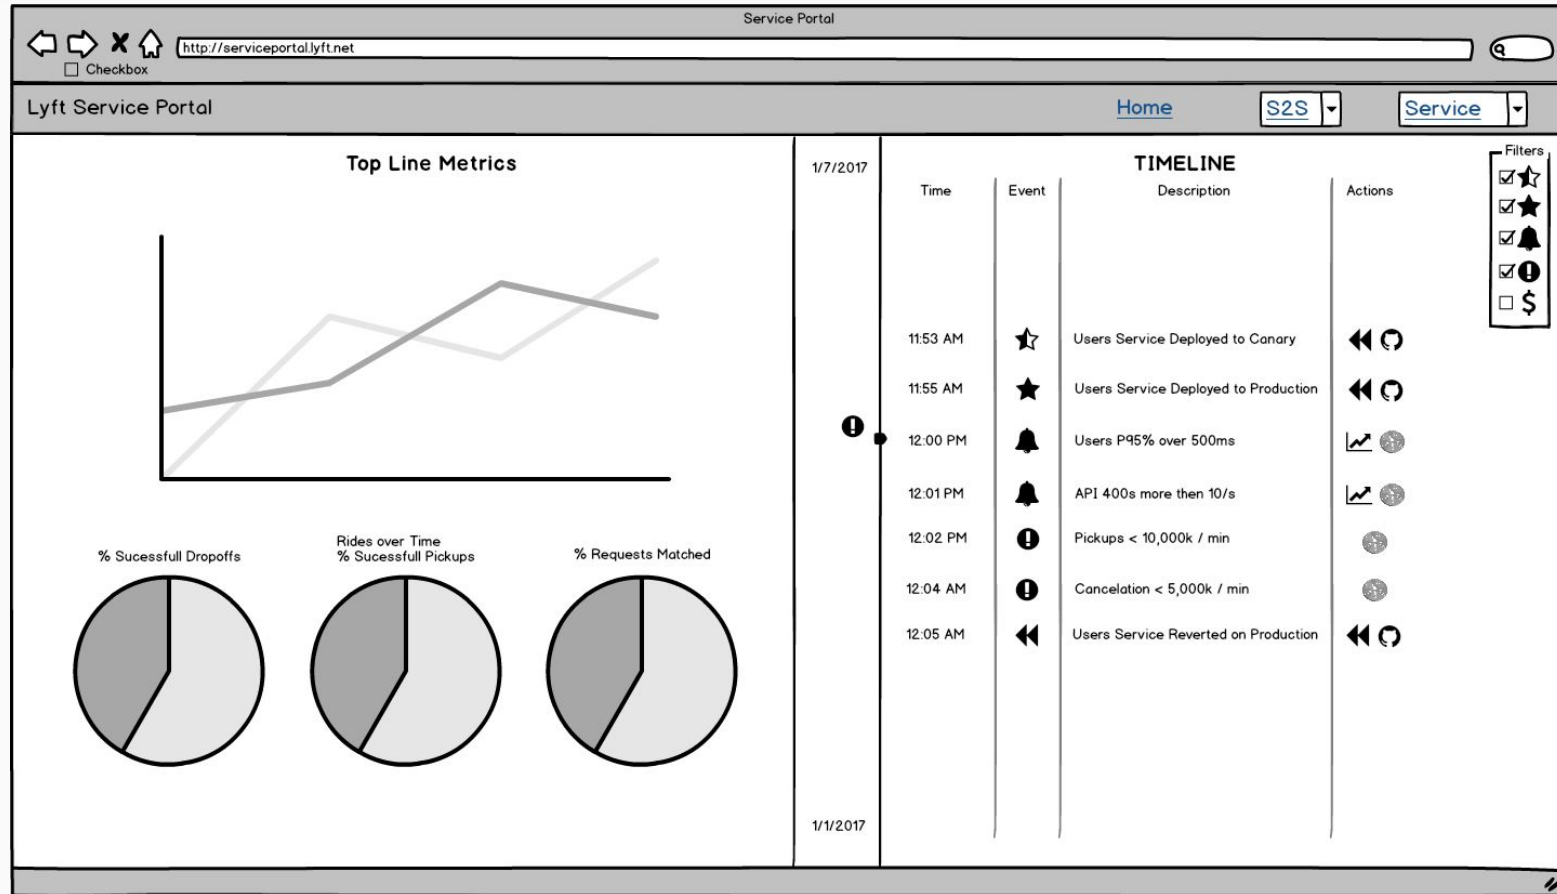
- Dev/Ops have too many data sources that are **not linked**.
- **Cognitive load** of different data sources make issue investigation with traditional stats, logging, and tracing is VERY high
- **Service mesh** yields an observability base that allows us to do *incredible things by default*.

How can we reimagine observability and operations in the age of the service mesh?

State of incident handling: Hystrix



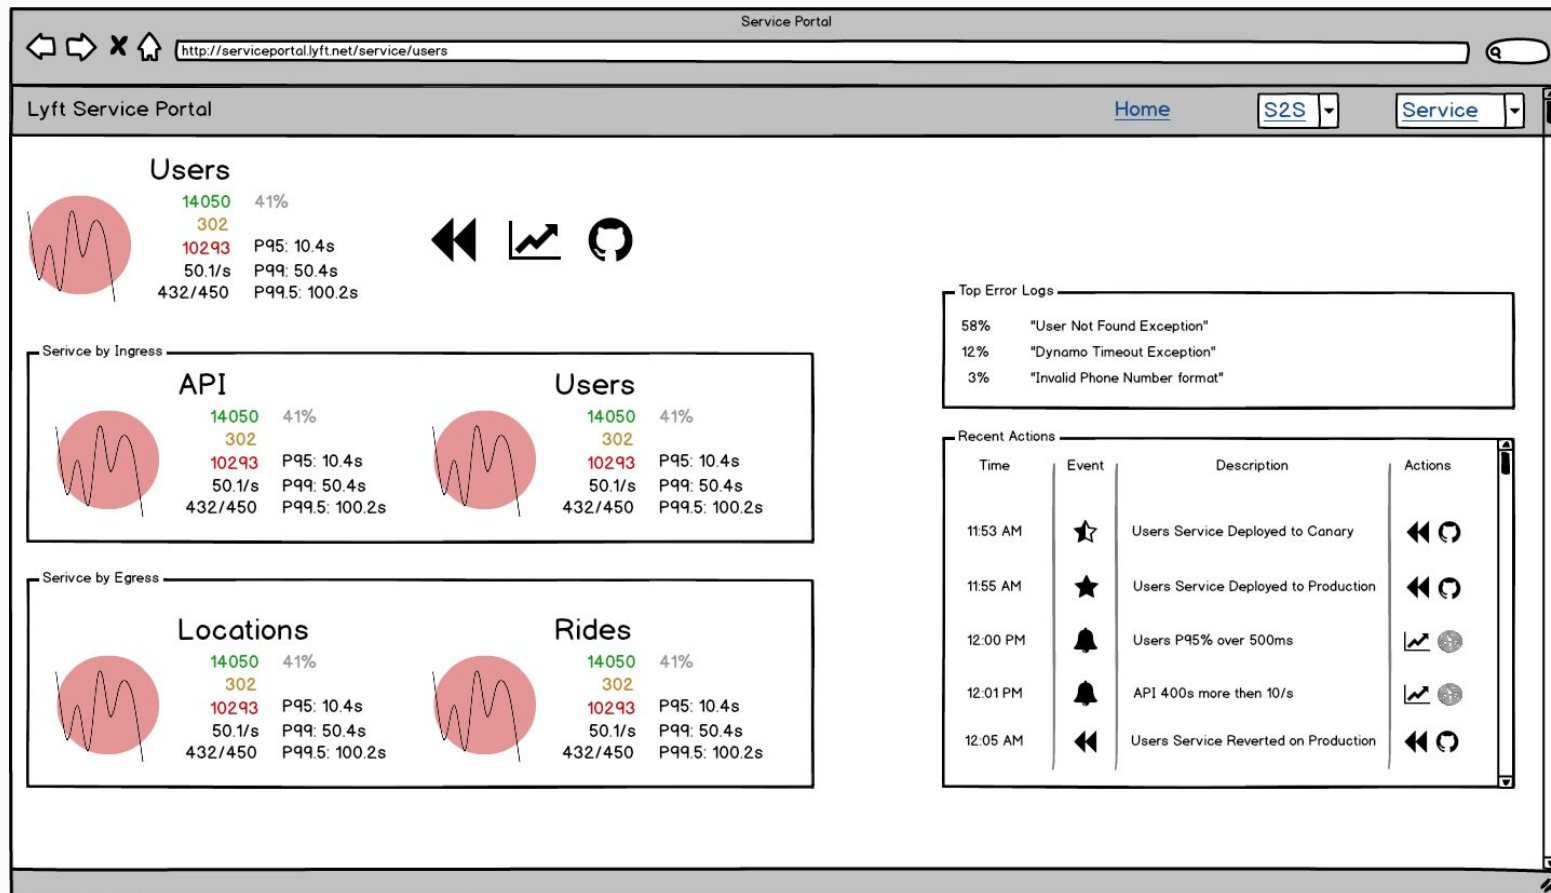
Service portal sketch: landing



Service portal sketch: service detail



Service portal sketch: service detail alternate





Optimal visualization of high level state



Actions relevant to mitigation



Machine learning to identify problems

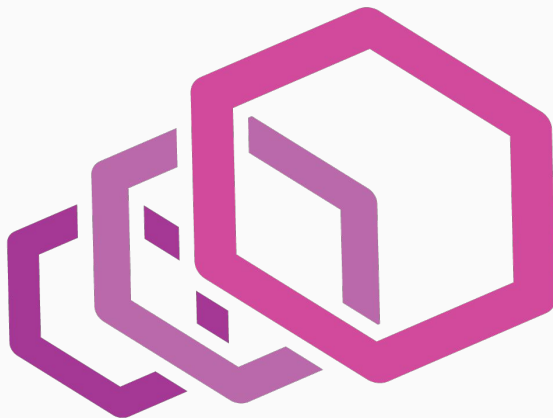


RBAC and versioning

How do we get there?

- A **universal data plane** like Envoy provides unified APIs for control as well as consistent observability output.
- Allows us to build **more feature-rich full service mesh solutions** such as Istio.
- When we assume the existence of the service mesh, we can **focus on an incredible UI/UX** instead of constantly trying to keep every application up to date.
- Assume that service mesh is the future... **All data is available.**
- We need to start building the UI/UX/ML of the future for distributed system command control. **Need to start now!**

- Thanks for coming! Questions welcome on Twitter: **@mattklein123**
- We are super excited about building a **community** around Envoy. Talk to us if you need help getting started.
- **Lyft is hiring!**



envoy