

GITTER CHANNEL

gitter.im/datawire/summit

TWEET ABOUT THE SUMMIT USING

#msvsummit

The Mechanics of Deploying Envoy at Lyft



Matt Klein

Senior Software Engineer at Lyft

Sponsored by




DATAWIRE



cloudbees






envoy
by **lyft**

The mechanics of deploying Envoy at Lyft
Microservices Practitioner Virtual Summit
Matt Klein / @mattklein123, Software Engineer @Lyft




Envoy refresher



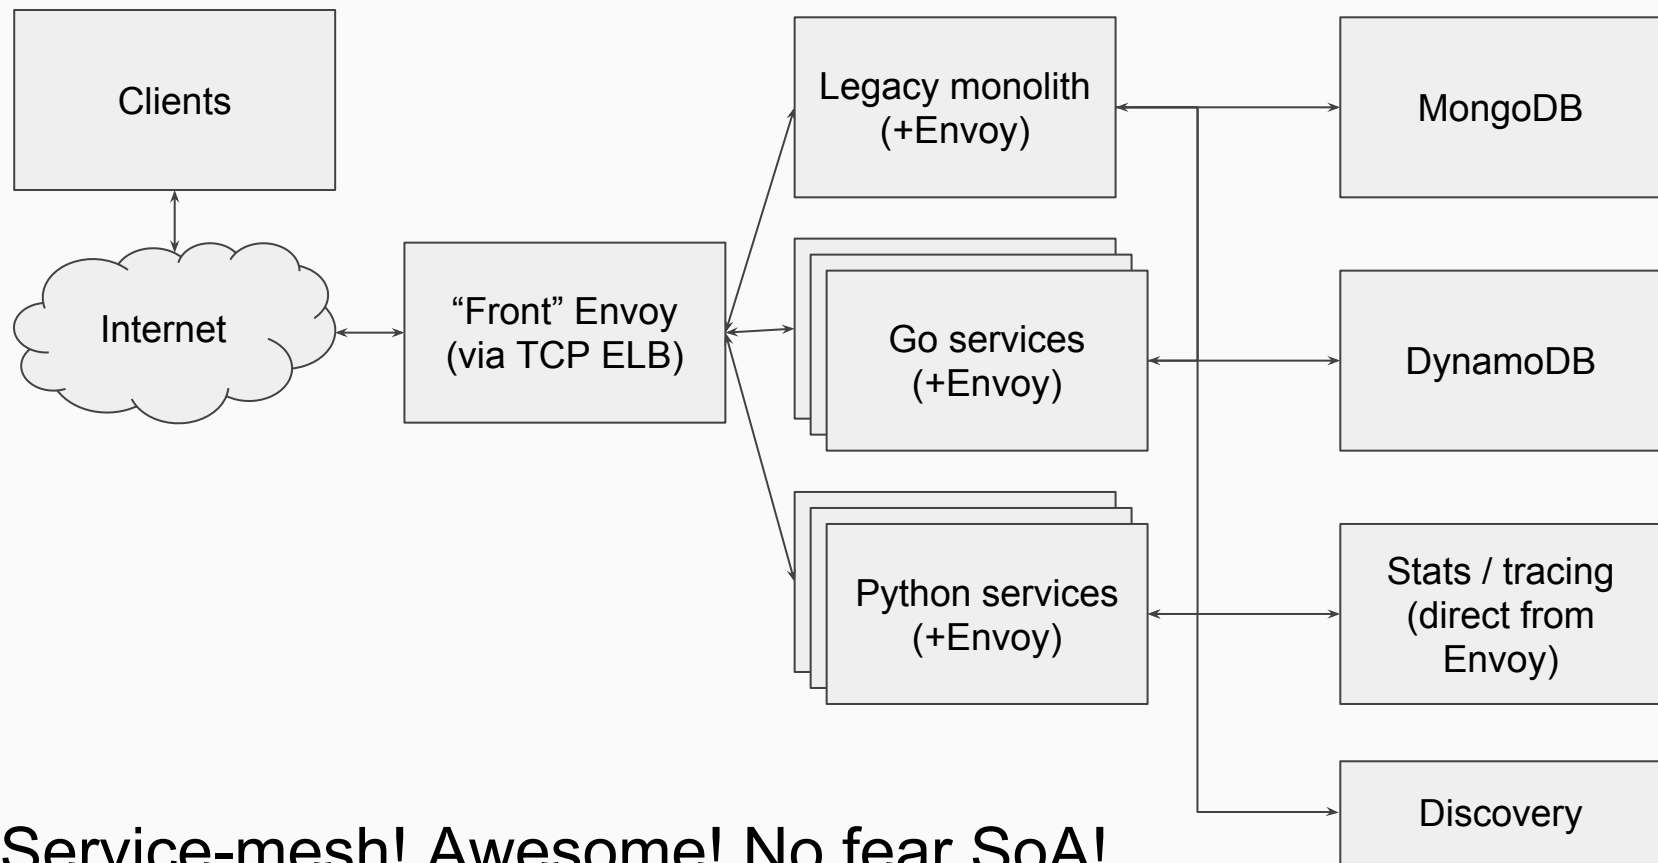
- **Out of process architecture:** Let's do a lot of really hard stuff in one place and allow application developers to focus on business logic.
- **Modern C++11 code base:** Fast and productive.
- **L3/L4 filter architecture:** A byte proxy at its core. Can be used for things other than HTTP (e.g., MongoDB, redis, stunnel replacement, TCP rate limiter, etc.).
- **HTTP L7 filter architecture:** Make it easy to plug in different functionality.
- **HTTP/2 first!** (Including **gRPC** and a nifty gRPC HTTP/1.1 bridge).
- **Service discovery and active/passive health checking.**
- **Advanced load balancing:** Retry, timeouts, circuit breaking, rate limiting, shadowing, outlier detection, etc.
- Best in class **observability:** stats, logging, and tracing.
- **Edge proxy:** routing and TLS.

Lyft ~4 years ago



Simple! No SoA! (*but still not that simple*)

Lyft today



Service-mesh! Awesome! No fear SoA!


How did Lyft go from then to now

- Can't go from before to after overnight.
- Must be incremental. Show value in steps.
- Perfect is the enemy of the good.
- This is how we did it...(teaser: it wasn't easy)



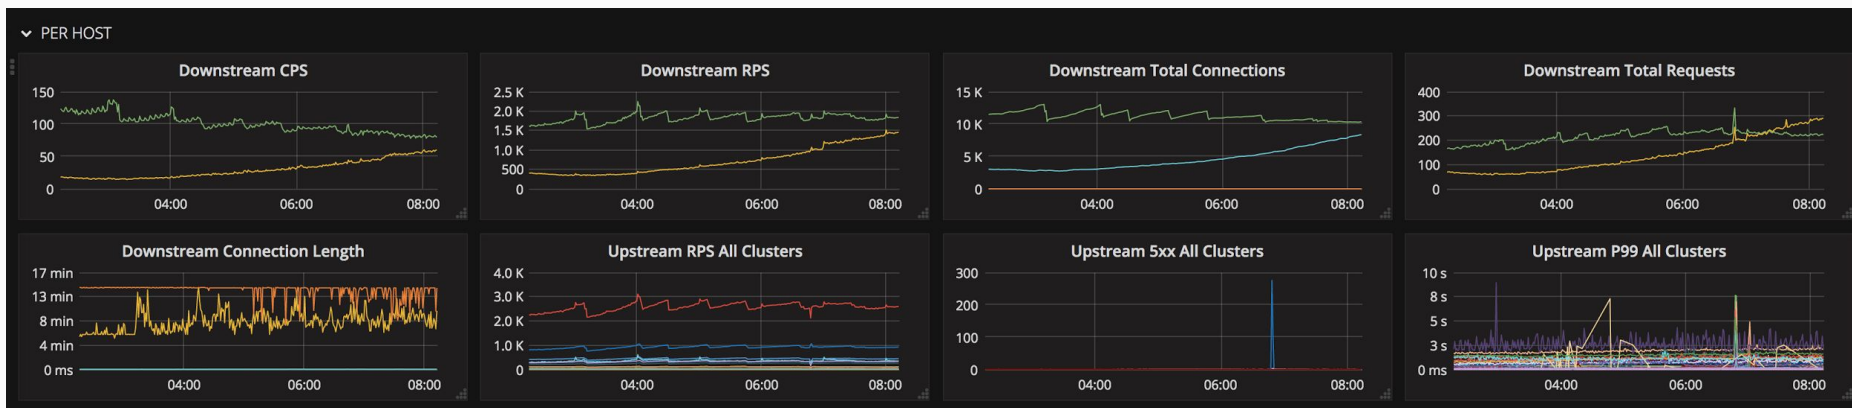
Start with edge proxy

- Microservice web apps need edge reverse proxy.
- Existing cloud offerings are not so good (even still).
- Easy to show value very quickly with stats, enhanced load balancing and routing, and protocols (h2/TLS).



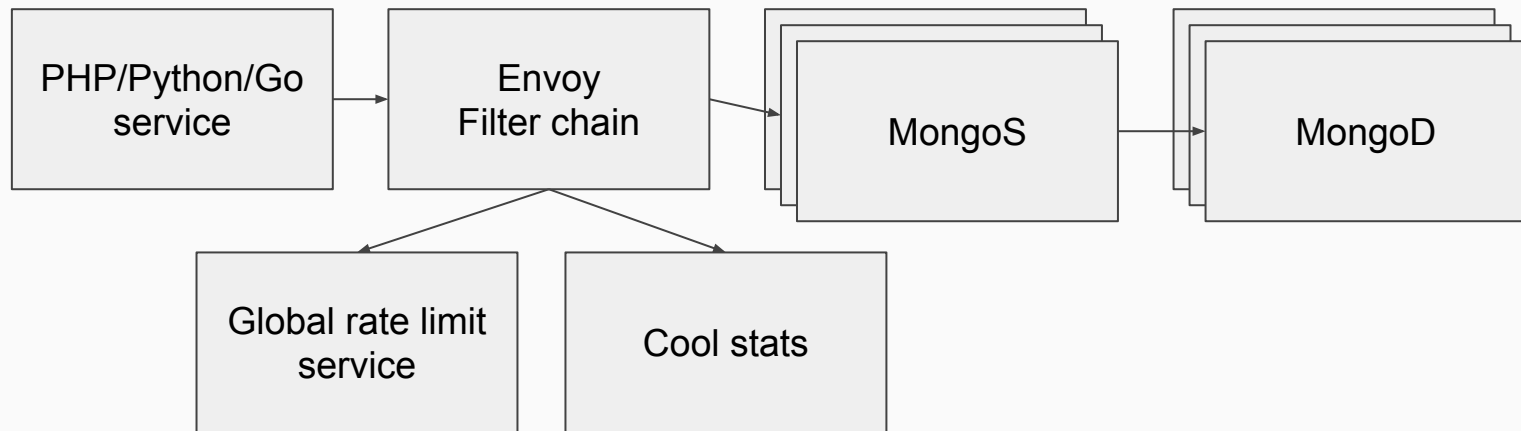
Start with edge proxy

Observability, observability, observability...




Next: TCP proxy and Mongo

- PHP to sharded Mongo not efficient with connection counts.
- Mongo bad at connection handling.
- Limit connections into Mongo.
- ... We can parse Mongo at L7 and generate cool stats!
- ... We can ratelimit Mongo to avoid death spirals!
- ... We can do this for all services, not just Python!




Next: service sidecar

- Because of Mongo, Envoy already running on services.
- Let's use for ingress buffering, circuit breaking, and observability.
- Still using internal ELBs at this point for service to service traffic.
- Still get tons of value out of above without direct mesh connect.



Next: service discovery and real mesh

- ELBs are (and intermediate LBs in general) ... not great. Let's do direct connect.
- Need service discovery.
- No ZK, etcd, or Consul. Eventually consistent. Build dead simple system with dead simple API.



Envoy thin clients @Lyft

```
from lyft.api_client import EnvoyClient
switchboard_client = EnvoyClient(
    service='switchboard'
)
msg = {'template': 'breaksignout'}
headers = {'x-lyft-user-id': 12345647363394}
switchboard_client.post("/v2/messages", data=msg, headers=headers)
```

- Abstract away egress port
- Request ID/tracing propagation
- Guide devs into good timeout, retry, etc. policies
- Similar thin clients for Go and PHP

Next ... Envoy everywhere

- OK this Envoy thing is pretty neat.
- Need to run it *everywhere* for full value.
- And so begins the slog. Many month burndown to convert everything and remove ELBs.
- Why a slog? Because of how Lyft manages services and configurations...
- But after slog complete, payoff is tremendous. Keep adding *features that developers want* and deploy them widely very quickly...



Initial Envoy config management

- Envoy config is JSON (please don't ask about JSON vs. YAML 🙄).
- Initially we had full JSON committed.
- Deploys compile Envoy and bundle configs. Don't need to do back compat.
- Config per deployment type.
 - front_envoy.json
 - service_to_service_envoy.json
- Envoy deployed via “pull deploy” and salt on each host.
- 1 service per host, 1 envoy per service, 1 envoy per host.
- Hot restart used for both config/binary deploys (no difference).

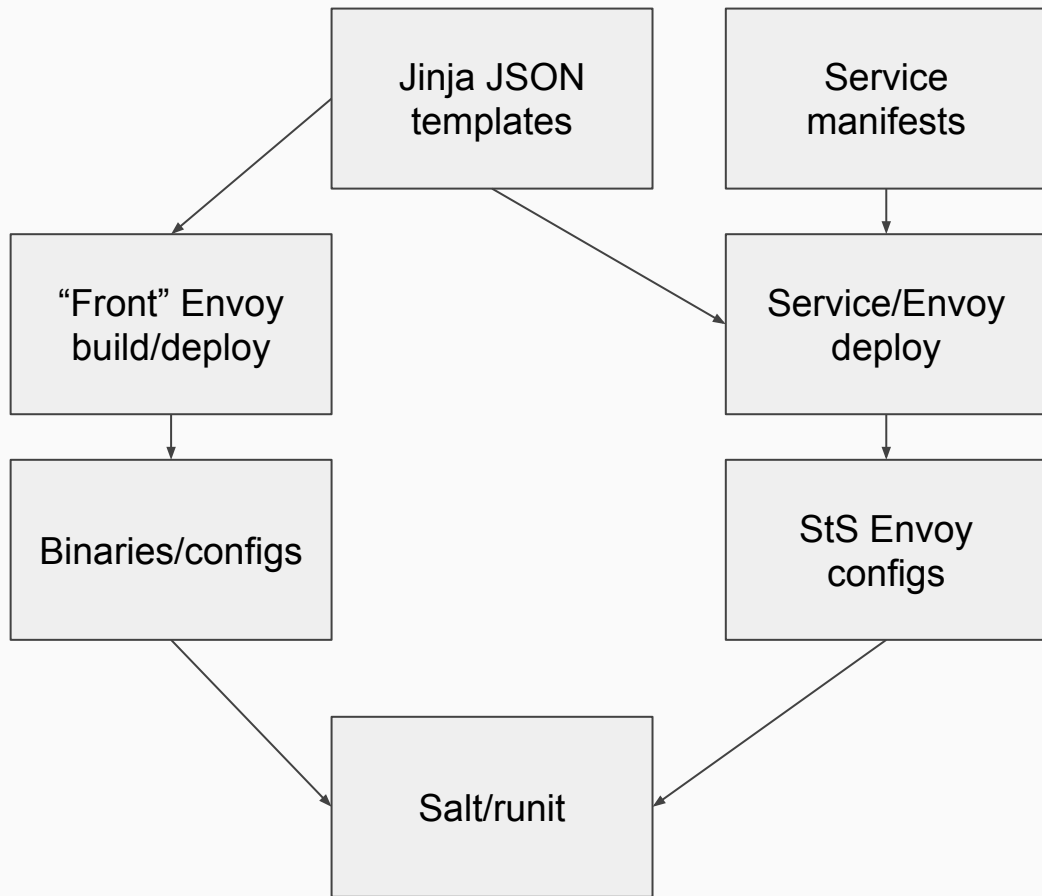
```
{  
  "listeners": {...},  
  "clusters": {...},  
  ...  
}
```

Next comes configgen.py and ... jinja!

- These configs are starting to get really tedious to write!
- Would like some amount of static scripting.
- Develop a tool called configgen.py which takes in a bunch of templates and inputs, runs jinja on them, and produces final outputs.
- Already starting to see trend that ultimately no way around complex Envoy configs being **machine generated**.

```
{% import 'certs.json' as certs -%}
{% macro listener(port,ssl,proxy_proto) %}
{
  "address": "tcp://0.0.0.0:{{ port }}",
  {% if ssl -%}
  "ssl_context": {
    "alpn_protocols": "h2,http/1.1",
    {{ certs.public(service_instance) }}
  },
  {% endif -%}
}
```


Envoy config/process management @Lyft now

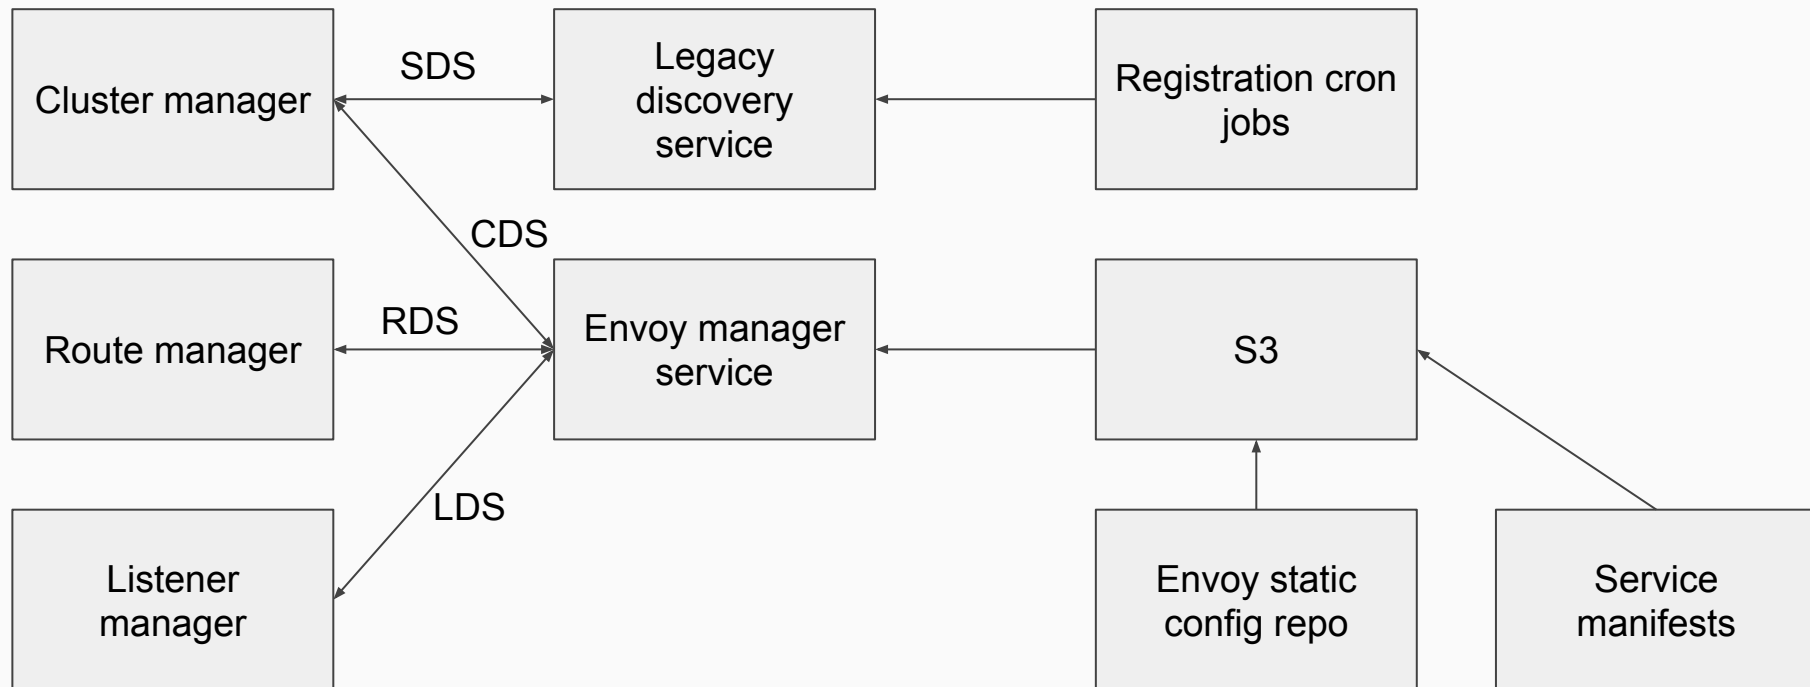


- Combination of static and dynamic configs.
- Service egress, circuit breaking, etc. configs specified in manifest.
- Service configs built on service host at service/envoy deploy time.

Envoy control plane APIs

- The goal of Envoy has always been as a *universal data plane*.
- Support APIs to enable control plane integration.
- V1 JSON/REST APIs (order of implementation):
 - **Service** Discovery Service (SDS) (note: poorly named)
 - **Cluster** Discovery Service (CDS)
 - **Route** Discovery Service (RDS)
 - **Listener** Discovery Service (LDS)
- @Lyft we **only currently use SDS**. Everything else driven by Jinja/JSON templating.

Lyft next-gen Envoy config management



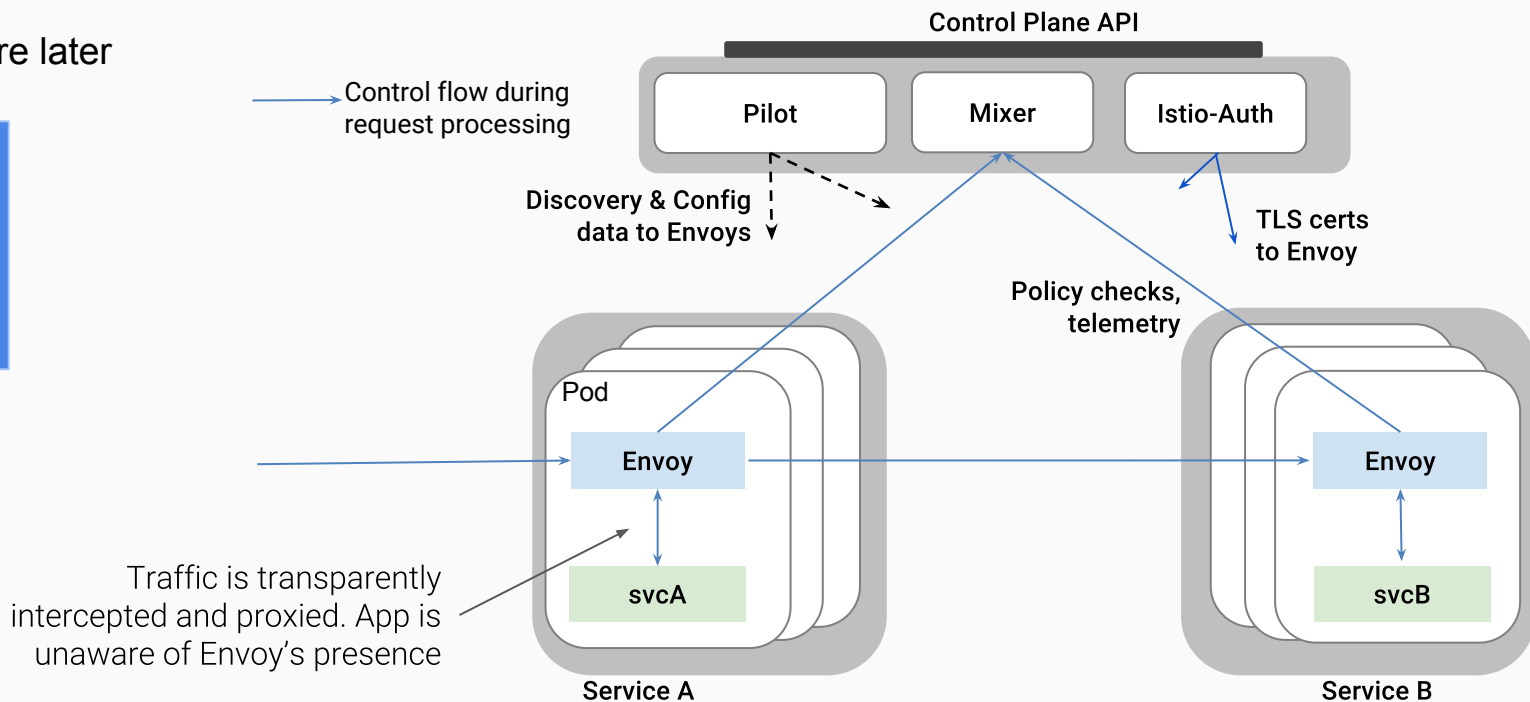
Only need a very tiny bootstrap config for each envoy...

Future gRPC/REST v2 APIs

- Canonical v2 APIs are **gRPC** (JSON also supported).
- Enhanced with **bidirectional streaming** and more features. Same spirit as v1.
- **Endpoint** Discovery Service (EDS, replaces SDS).
 - Fetch endpoints/hosts.
 - Report load info.
- **Cluster** Discovery Service (CDS).
- **Route** Discovery Service (RDS).
- **Listener** Discovery Service (LDS).
- **Health** Discovery Service (HDS).
- **Aggregated** Discovery Service (ADS).
 - Allowing all xDS APIs to be served by a single stream if desired. Useful for enforcing ordering of updates. E.g., CDS -> EDS -> RDS.

Istio and config APIs

- “Raw Envoy” is still too hard to configure for most folks
- Hard to decouple network from deploy orchestration
- Build higher level control systems (universal dataplane)
- Istio!
- K8s first. More later



- Thanks for coming! Questions welcome on Twitter: @mattklein123
- We are super excited about building a **community** around Envoy. Talk to us if you need help getting started.
- <https://lyft.github.io/envoy/> (@envoyproxy)
- <https://istio.io/> (@istiomesh)

