

Kubernetes Philosophy

Configuration Management Best Practice

Mengqi Yu, Software Engineer, Google

Kubernetes Concepts

Kubectl Commands

Configuration Management

Demo

Declarative App Management

主办:



Kubernetes Concepts

- Declarative
 - But kubernetes support imperative
 - Support self-healing
- Level-triggered
- Asynchronous

Declarative vs Imperative

Declarative:

Definition: user specifies what the desired state is.
The system will know what to do to move current state to desired state.

user: set desired state to 3



observed: 2



system: +1



observed: 3



Imperative:

Definition: user tells the system what to do.

user: +1



system: +1



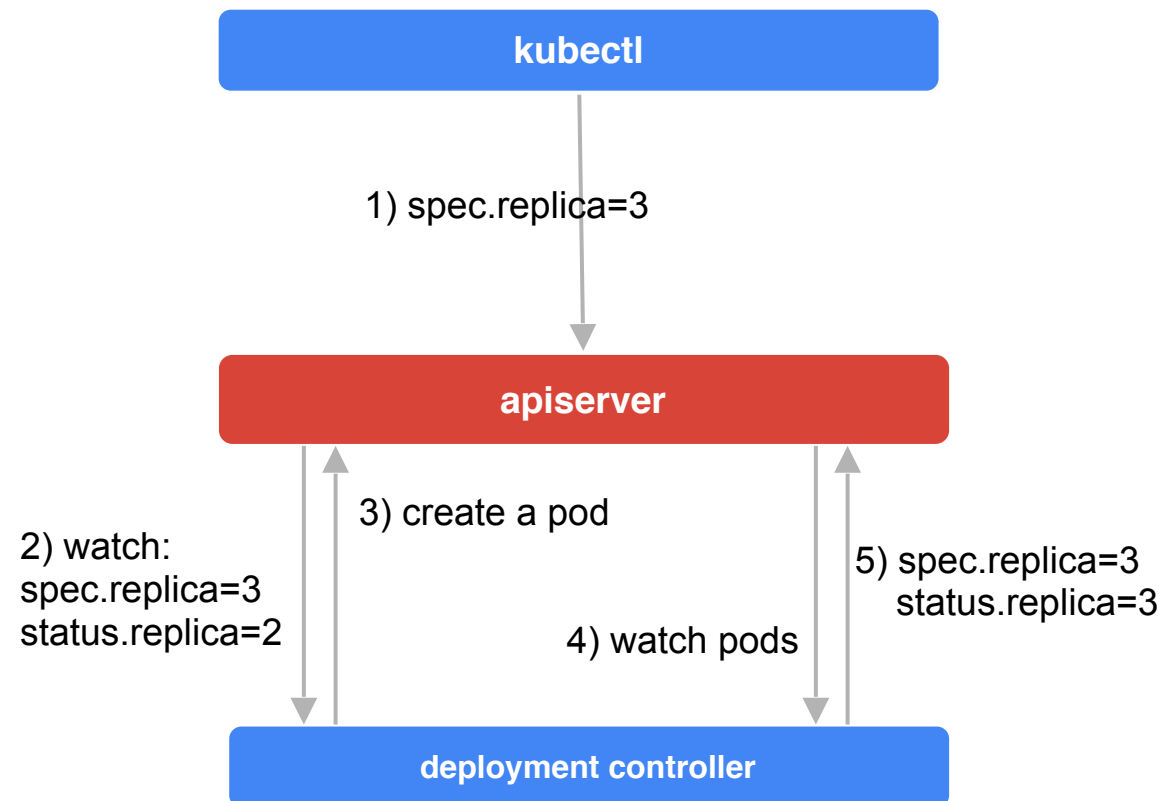
Spec vs Status

Most kubernetes API objects have Spec and Status fields.

- Spec: user desired state
 - Used by reconciliation loops to update other objects in the cluster or external objects.
- Status: system observed state
 - Used by clients and reconciliation loops to view the state of the object.

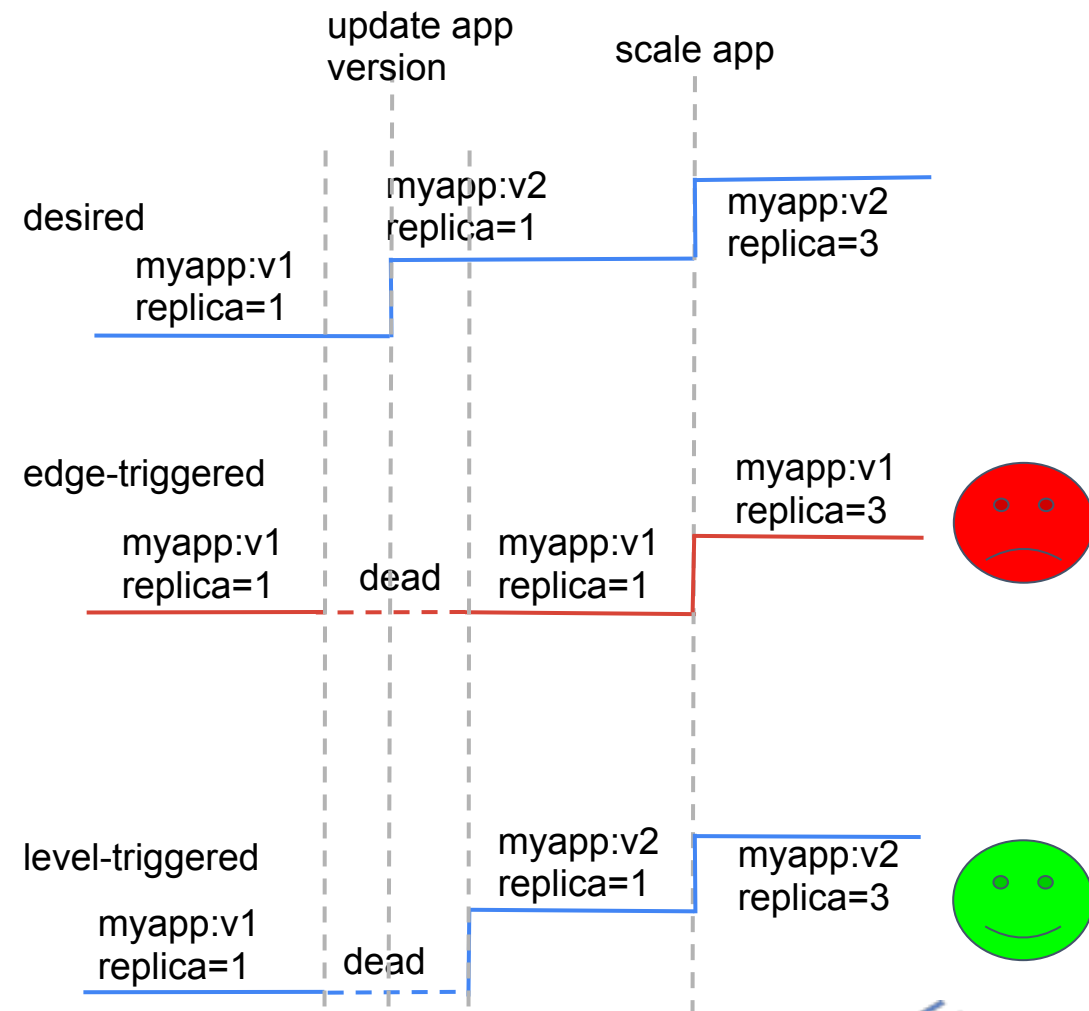
Reconcile

- Controller watches the resources.
- Desired state changes.
- Controller is notified of an update and compares current desired vs current observed states.
- Controller responds to actual-desired mismatch by executing operations (e.g. create a Pod, allocate clusterIp).
- Observed state matches desired state.



Level-triggered vs Edge-triggered

- Kubernetes is level-triggered
- Edge-triggered
 - Pros: straightforward and easy to implement.
 - Cons:
 - missing edges may result in wrong states.
 - does not take most direct route to current desired state
- Level-triggered
 - Pros: don't need to worry about missing the edge.
 - Cons: harder to implement.



主办:



Kubernetes Concepts

Kubectl Commands

Configuration Management

Demo

Declarative App Management

主办：



Imperative Commands

- **Imperative commands:** kubectl supports imperative create / update / delete commands that parse command line flags into yaml / json to send to the server.
 - This hides the underlying object configuration for the users.
 - e.g. kubectl run, expose
 - e.g. kubectl create deployment
 - e.g. kubectl delete <resource_kind>/<resource_name>

Imperative Object Configuration

- **Object configuration:** objects are defined in yaml / json with syntax defined in reference docs.
- **Imperative object configuration:** kubectl supports imperative create / update / delete commands that work directly on **object configuration** stored locally or in scm.
 - e.g. kubectl create -f
 - e.g. kubectl replace -f
 - e.g. kubectl delete -f

Imperative Commands vs Imperative Object Configuration



- Imperative commands
 - Pros: straightforward, easy to learn and easy to remember.
 - Cons:
 - no change review processes.
 - no audit trails.
 - not provide a template.
- Imperative object configuration
 - Pros:
 - integrate with a source control system, e.g. git.
 - integrate with reviewing changes before push processes.
 - integrate with audit trails.
 - Cons:
 - requires basic understanding of the object schema.
 - need to write the yaml files.

Declarative Object Configuration

- **Declarative object configuration:** kubectl supports declarative workflow that works directly on **object configuration**.
 - User does not specify create / update / delete.
 - Command figures out the operation.
 - e.g. kubectl apply --prune -f (a.k.a configuration reconciliation)

- Imperative object configuration
 - Pros: simpler and easier to understand.
 - Cons:
 - best on files, not directories.
 - updates by other writers (e.g. controller) MUST be reflected in configuration files.
- Declarative object configuration
 - Pros:
 - updates to live objects by other writers are retained. - e.g. autoscalers change replicas
 - works with directories
 - automatically choose operations (create, patch, delete)
 - Cons:
 - behavior is harder to understand

Kubernetes Concepts

Kubectl Commands

Configuration Management

Demo

Declarative App Management

主办：

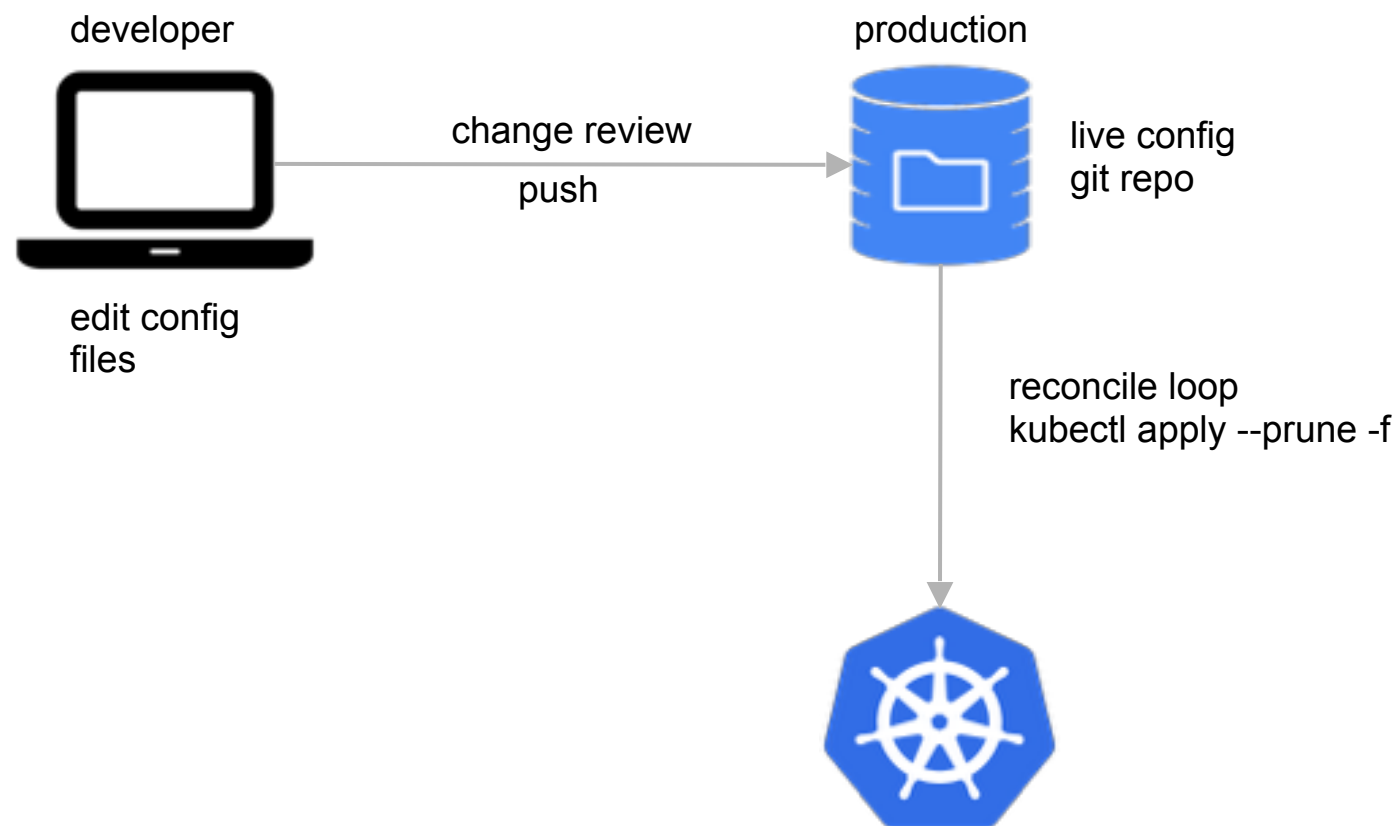


Configuration Management

- Beginning user
 - Start with imperative commands
 - Work with imperative object configuration
- Experienced User
 - Use declarative object configuration
 - Use version control
 - Use change review process

- Beginning user
 - Kubernetes API is the only store for desired state.
 - Support some number of revision history
 - But no long history, no cross-object history
- Experienced User
 - Revision control contains user defined desired state.
 - Support cross-object history.
 - Multiple deployments that depend on each other: tag the whole repository of cluster state at known good points, can roll back entire cluster by applying state at that tag.
 - Auto reconciliation e.g. [kube-applier](#)

Workflow



Demo

主办：



Kubernetes Concepts

Kubectl Commands

Configuration Management


Demo

Declarative App Management

主办:



Declarative Application Management Workflow



Declarative Application Management (DAM)

- Lookup the off-the-shelf application in an app registry
- Fork
- (Optional) Make change of it: labels, name prefix
- (Optional) Create multiple instances of overlays (e.g. dev, staging, prod)
- Check them in your git repo
- CD system auto reconciliation

DAM is actively being explored but in early development phase.

- Kubernetes Charts
 - Curated applications for Kubernetes.
 - <http://github.com/kubernetes/charts>
- Kubernetes Helm
 - A tool for manage kubernetes packages.
 - <https://github.com/kubernetes/helm>
- Challenge: it's hard to find the balance between enough parameters to satisfy all users, without the number of parameters becoming hard to maintain and understand
 - e.g. [Jenkins chart](#)

Recap

- Declarative vs Imperative
- Spec vs Status
- Kubectl commands
- Configuration management
 - Version control system

Q&A

Thank you for your time

主办：

