

Kubernetes 存储技术介绍及在云计算平台 Azure 上的实现

张夏 微软上海开源技术组

Github: andyzhangx

Kubernetes持久化存储(PV)技术介绍

k8s持久化存储(PV)在Azure上的实现

如何实现持久化存储(PV)插件

主办：



PersistentVolume(PV)

提供网络存储资源 by admin

PersistentVolumeClaim (PVC)

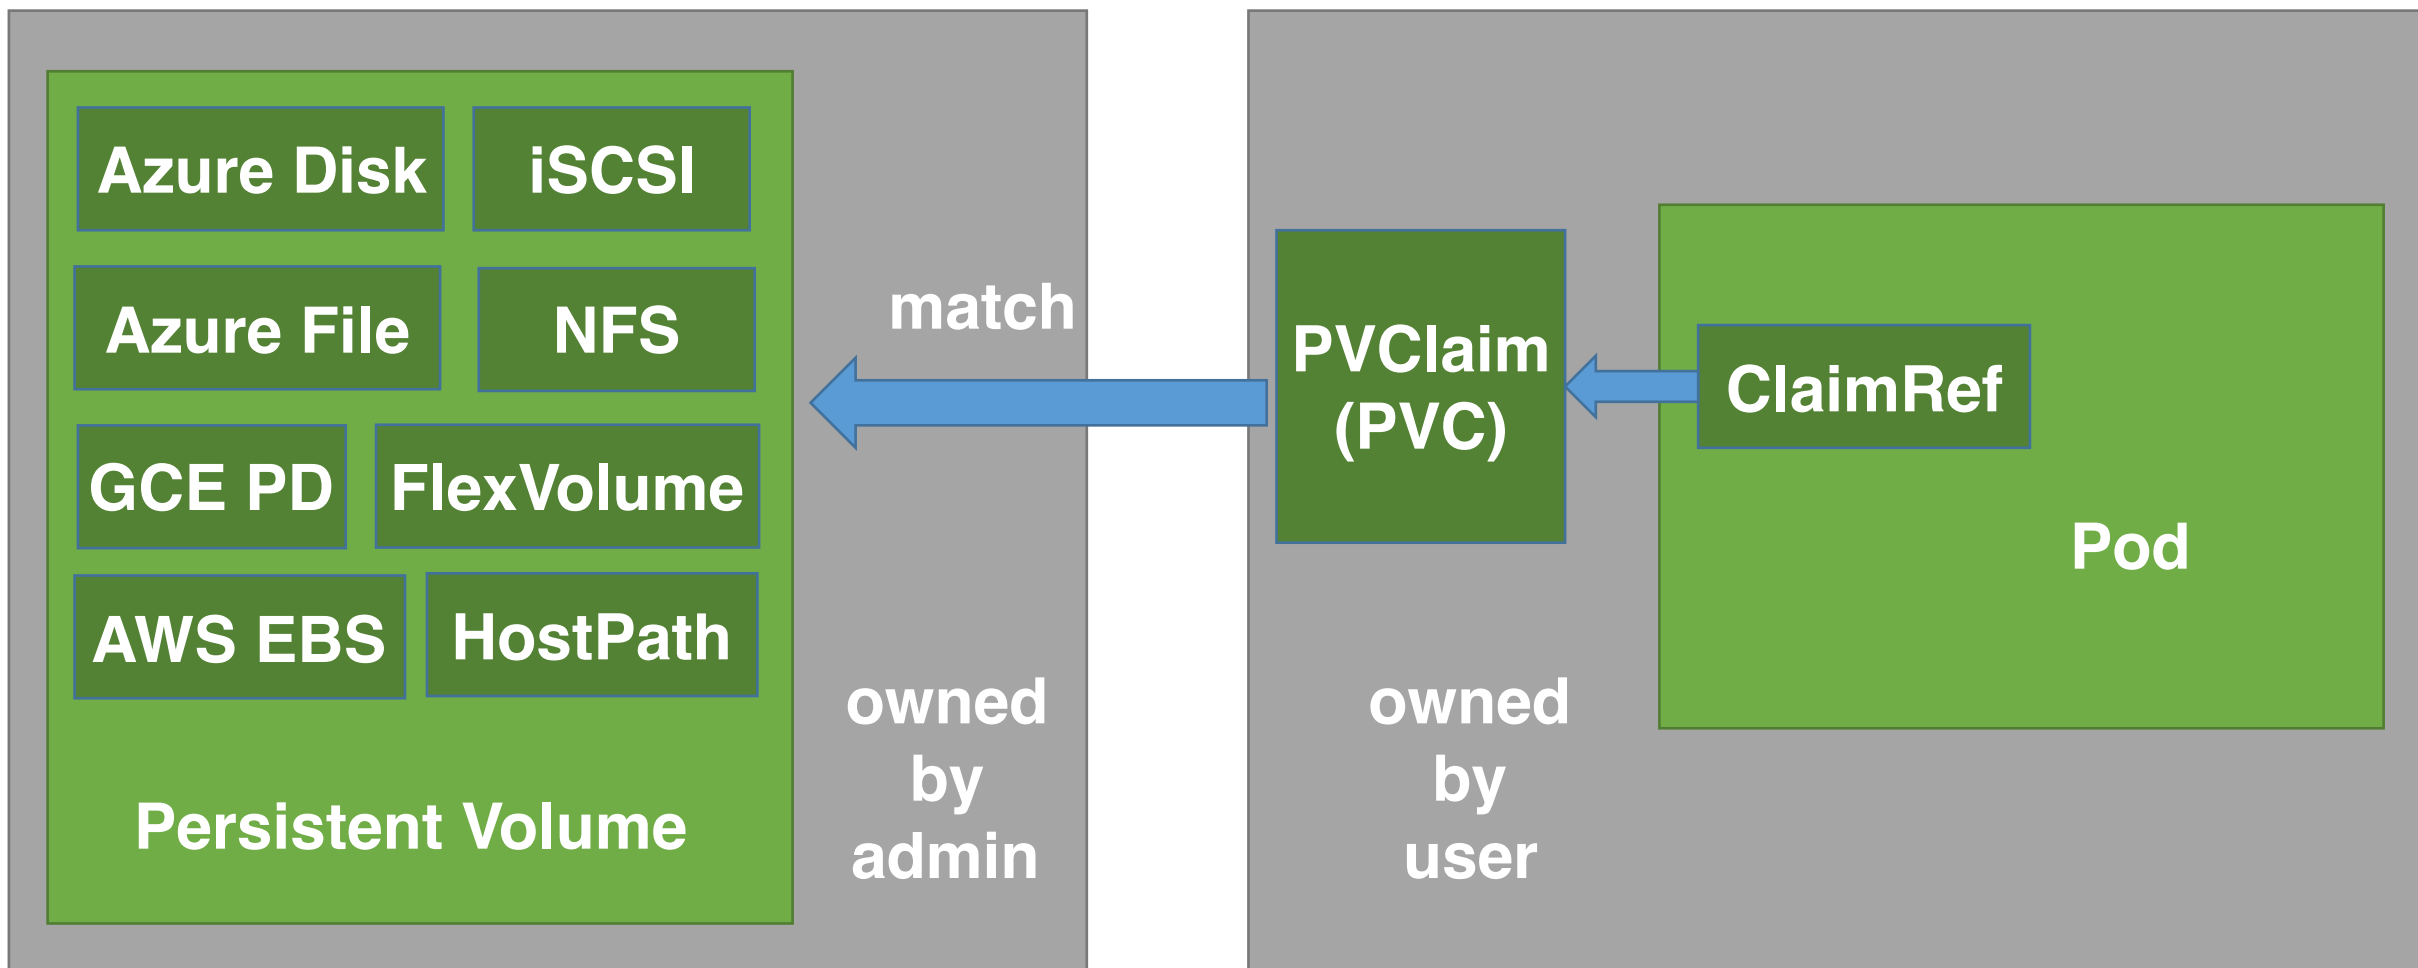
请求存储资源 by user

StorageClass

封装不同类型存储 by admin

动态创建PV

PV & PVC



Spec

Capacity 容量

storage size

Access Modes 访问模式

RWO, RWX, ROX

Class 存储类

Reclaim Policy 回收策略

Retain, Recycle, Delete

Mount Options 加载选项

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```

Status

Available, Bound, Released, Failed

主办:



Spec

Access Modes 访问模式

RWO, RWX, ROX

Resources 请求资源

storage size

Class 存储类

Selector 匹配

match label

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 8Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}
```

Status

Pending, Bound

主办:



Provisioner

Internal(In-Tree) Provisioner

External(Out-of-Tree) Provisioner

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: slow
provisioner: kubernetes.io/azure-disk
parameters:
  storageaccounttype: Standard_LRS
  kind: Shared
```

Reclaim Policy 回收策略

Parameters 参数

主办:



Volume Plugin	Internal Provisioner	Config Example
AWSElasticBlockStore	✓	AWS
AzureFile	✓	Azure File
AzureDisk	✓	Azure Disk
CephFS	-	-
Cinder	✓	OpenStack Cinder
FC	-	-
FlexVolume	-	-
Flocker	✓	-
GCEPersistentDisk	✓	GCE
Glusterfs	✓	Glusterfs
iSCSI	-	-
PhotonPersistentDisk	✓	-
Quobyte	✓	Quobyte
NFS	-	-
RBD	✓	Ceph RBD
VsphereVolume	✓	vSphere
PortworxVolume	✓	Portworx Volume

Kubernetes持久化存储(PV)技术介绍

k8s持久化存储(PV)在Azure上的实现

如何实现持久化存储(PV)插件

主办：



CLOUD NATIVE
COMPUTING FOUNDATION



AzureDisk

通过在Node上加载Azure磁盘方式实现

Access Mode:

ReadWriteOnce(RWO)

AzureFile

通过在Node上加载Azure共享文件（SMB 2.1, 3.0 ）实现

Access Mode:

ReadWriteOnce(RWO), ReadWriteMany(RWX), ReadOnlyMany(ROX)

1. 创建Azure Disk StorageClass
2. 创建PersistentVolumeClaim(PVC)
3. 创建Pod时使用该PVC

```
kind: StorageClass
apiVersion: v1
metadata:
  name: hdd
provisioner: kubernetes.io/azure-disk
parameters:
  storageaccounttype: Standard_LRS
  kind: Shared
```

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: pvc-azuredisk
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  storageClassName: hdd
```

使用Azure Disk Volume

1. `kubectl create -f storageclass-azuredisk.yaml`
2. `kubectl create -f pvc-azuredisk.yaml`

```
Every 2.0s: kubectl describe pvc pvc-azuredisk Sat Sep 23 02:18:55 2017

Name:          pvc-azuredisk
Namespace:     default
StorageClass:  hdd
Status:        Bound
Volume:        pvc-600956a1-a005-11e7-a4df-000d3a36d70c
Labels:        <none>
Annotations:   pv.kubernetes.io/bind-completed=yes
               pv.kubernetes.io/bound-by-controller=yes
               volume.beta.kubernetes.io/storage-provisioner=kubernetes.io/azure-disk
Capacity:      5Gi
Access Modes:  RWO
Events:
  FirstSeen    LastSeen    Count   From              SubObjectPath  Type    Reason
  ----
  32s          32s         1      persistentvolume-controller  Normal        ProvisioningSucceeded
Successfully provisioned volume pvc-600956a1-a005-11e7-a4df-000d3a36d70c using kubernetes.io/azure-disk
```


创建一个基于Azure Disk Volume的Windows POD

```
kind: Pod
apiVersion: v1
metadata:
  name: aspnet-azuredisk
spec:
  containers:
    - image: microsoft/aspnet
      name: aspnet-azuredisk
      volumeMounts:
        - name: blobdisk01
          mountPath: 'D:'
  nodeSelector:
    beta.kubernetes.io/os: windows
  volumes:
    - name: blobdisk01
      persistentVolumeClaim:
        claimName: pvc-azuredisk
```



Every 2.0s: kubectl describe po aspnet-azuredisk

Sat Sep 23 02:45:13 2017

Name: aspnet-azuredisk
Namespace: default
Node: 14777acs9000/10.240.0.5
Start Time: Sat, 23 Sep 2017 02:43:37 +0000
Labels: name=storage
Annotations: <none>
Status: Running
IP: 10.244.1.25

Containers:

aspnet-azuredisk:
Container ID: docker://1df89ab285ef805b16c72671ceefd71a2332408da5e3f875775ac9b008539108
Image: microsoft/aspnet
Image ID: docker-pullable://microsoft/aspnet@sha256:3456f41d5745c1992a82e66afcfb16e6c203c4af5ae3013f4b26c897683748a2
Port: <none>
State: Running
Started: Sat, 23 Sep 2017 02:44:57 +0000
Ready: True
Restart Count: 0
Environment: <none>

Mounts:

/var/run/secrets/kubernetes.io/serviceaccount from default-token-7gh5d (ro)
D: from blobdisk01 (rw)

Conditions:

Type	Status
Initialized	True
Ready	True
PodScheduled	True

Volumes:

blobdisk01:
Type: PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
ClaimName: pvc-azuredisk
ReadOnly: false
default-token-7gh5d:
Type: Secret (a volume populated by a Secret)
SecretName: default-token-7gh5d
Optional: false

QoS Class: BestEffort
Node-Selectors: beta.kubernetes.io/os=windows

Tolerations: <none>

Events:

FirstSeen	LastSeen	Count	From	SubObjectPath	Type	Reason	Message
1m	1m	1	default-scheduler		Normal	Scheduled	Successfully assigned aspnet-azuredisk to 14777acs9000
1m	1m	1	kubelet, 14777acs9000		Normal	SuccessfulMountVolume	MountVolume.SetUp succeeded for volume "default-token-7gh5d"
23s	23s	1	kubelet, 14777acs9000		Normal	SuccessfulMountVolume	MountVolume.SetUp succeeded for volume "pvc-600956a1"

主办:



使用Azure Disk Volume

```
azureuser@k8s-master-14777338-0:~$ kubectl exec -i aspnet-azuredisk -- cmd
Microsoft Windows [Version 10.0.14393]
(c) 2016 Microsoft Corporation. All rights reserved.

C:\>d:

D:\>dir
Volume in drive D has no label.
Volume Serial Number is 56EB-6E22

Directory of D:\

09/25/2017  02:31 AM    <DIR>          test
               0 File(s)                0 bytes
               1 Dir(s)          5,334,323,200 bytes free
```

On Kubernetes Master (Controller Manager)

- 在k8s 集群所在的资源组中创建一块Azure Disk
- 将新创建的Azure Disk添加到kubernetes node上

On Kubernetes Node (Kubelet)

- 循环检测新磁盘
- 将检测到的新磁盘格式化并链接到kubernetes目录下
- 创建pod时使用volume映射，将该目录映射到用户指定目录

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: azurefile
provisioner: kubernetes.io/azure-file
parameters:
  skuName: Standard_LRS
  location: westus
```

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: pvc-azurefile
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 5Gi
  storageClassName: azurefile
```


使用Azure File Volume

1. `kubectl create -f storageclass-azurefile.yaml`
2. `kubectl create -f pvc-azurefile.yaml`

```
Every 2.0s: kubectl describe pvc pvc-azurefile Sat Sep 23 02:35:29 2017

Name:          pvc-azurefile
Namespace:     default
StorageClass:  azurefile
Status:        Bound
Volume:        pvc-af592494-a007-11e7-8139-000d3a312c80
Labels:        <none>
Annotations:   pv.kubernetes.io/bind-completed=yes
               pv.kubernetes.io/bound-by-controller=yes
               volume.beta.kubernetes.io/storage-provisioner=kubernetes.io/azure-file
Capacity:      5Gi
Access Modes:  RWX
Events:
  FirstSeen    LastSeen    Count   From              SubObjectPath  Type    Reason
  ----
  1m           1m          1      persistentvolume-controller  Normal    ProvisioningSucceeded
  Successfully provisioned volume pvc-af592494-a007-11e7-8139-000d3a312c80 using kubernetes.io/azure-file
```

创建一个基于Azure File Volume 的 pod

```
kind: Pod
apiVersion: v1
metadata:
  name: nginx-azurefile
spec:
  containers:
    - image: nginx
      name: nginx-azurefile
  volumeMounts:
    - name: azurefile01
      mountPath: "/mnt/blobfile"
  volumes:
    - name: azurefile01
      persistentVolumeClaim:
        claimName: pvc-azurefile
```



Every 2.0s: kubectl describe po nginx-azurefile

```

Name:          nginx-azurefile
Namespace:     default
Node:          k8s-agent-c3ee7eel-1/10.240.0.4
Start Time:    Sat, 30 Sep 2017 14:30:04 +0000
Labels:        <none>
Annotations:   <none>
Status:        Running
IP:            10.244.0.8
Controllers:   <none>
Containers:
  nginx-azurefile:
    Container ID:   docker://1b95873aa8add07ddba83f929b52f5819ffddfae8818378e626975df7948c7c
    Image:          nginx
    Image ID:       docker-pullable://nginx@sha256:af32e714a9cc3157157374e68c818b05ebe9e0737aac06b55a09da374209a8f9
    Port:
    Command:
      /bin/sh
      -c
      while true; do echo $(date) >> /mnt/blobfile/outfile; sleep 1; done
    State:          Running
      Started:      Sat, 30 Sep 2017 14:30:09 +0000
    Ready:          True
    Restart Count:   0
    Environment:    <none>
    Mounts:
      /mnt/blobfile from blobfile01 (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-zjth4 (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready          True
  PodScheduled    True
Volumes:
  blobfile01:
    Type:          PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
    ClaimName:     pvc-azurefile
    ReadOnly:      false
  default-token-zjth4:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-zjth4
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     <none>
Events:
  FirstSeen    LastSeen    Count   From              SubObjectPath               Type            Reason            Message
  -----
  24s          24s         1      default-scheduler                Normal          Scheduled         Successfully assigned nginx-azurefile to k8s-
agent-c3ee7eel-1
  24s          24s         1      kubelet, k8s-agent-c3ee7eel-1    MountVolume     SuccessfulMountVolume MountVolume.Setup succeeded for volume "defa

```

工外

使用Azure File Volume

```
azureuser@k8s-master-C3EE7EE1-0:~$ kubectl exec -it nginx-azurefile -- bash
root@nginx-azurefile:/# df -h
Filesystem                                Size  Used Avail Use% Mounted on
overlay                                  30G   8.4G   21G   29% /
tmpfs                                    3.5G     0   3.5G    0% /dev
tmpfs                                    3.5G     0   3.5G    0% /sys/fs/cgroup
/dev/sda1                                30G   8.4G   21G   29% /etc/hosts
//andytestd.file.core.windows.net/andyk8s162-dynamic-pvc-da52587b-a386-11e7-9229-000d3a35d0d5  5.0G   64K   5.0G    1% /mnt/blobfile
shm                                       64M     0   64M    0% /dev/shm
tmpfs                                    3.5G   12K   3.5G    1% /run/secrets/kubernetes.io/serviceaccount
```

On Kubernetes **Master (Controller Manager)**

- 在k8s cluster所在的资源组中创建Azure File(SMB share)

On Kubernetes **Node (Kubelet)**

- 取得azure file账号信息
- 将创建的Azure File 使用cifs协议加载到kubernetes目录下

Azure Disk & File Volume Demo

<https://github.com/andyzhangx/Demo/tree/master/windows/azuredisk>

<https://github.com/andyzhangx/Demo/tree/master/linux/azurefile>

Kubernetes持久化存储(PV)技术介绍

k8s持久化存储(PV)在Azure上的实现

如何实现持久化存储(PV)插件

主办：



实现新的Internal Cloud Provider?

kubernetes / kubernetes

Watch

1,954

Star

27,306

Fork

9,787

Code

Issues 4,597

Pull requests 871

Projects 11

Wiki

Insights

Branch: master

kubernetes / pkg / cloudprovider / providers /

Create new file

Upload files

Find file

History

k8s-merge-robot Merge pull request #53276 from freehan/alpha-backendservice

Latest commit 0ca0d76 2 days ago

aws	Use custom error for "unimplemented"	6 days ago
azure	Use custom error for "unimplemented"	6 days ago
cloudstack	godep: add dhcp4 and dhcp4client dependencies	13 days ago
fake	Use custom error for "unimplemented"	6 days ago
gce	Merge pull request #53276 from freehan/alpha-backendservice	2 days ago
openstack	Merge pull request #53101 from FengyunPan/default-external	3 days ago
ovirt	Use custom error for "unimplemented"	6 days ago
photon	Merge pull request #52983 from kubermatic/use-notimplemented-error	3 days ago
vsphere	Merge pull request #52983 from kubermatic/use-notimplemented-error	3 days ago
BUILD	remove rackspace related code	10 days ago
providers.go	remove rackspace related code	10 days ago

实现新的Internal Cloud Provider?

kubernetes / kubernetes

Watch 1,945 Star 27,180 Fork 9,734

Code Issues 4,588 Pull requests 871 Projects 11 Wiki Insights

Branch: master kubernetes / pkg / cloudprovider /

Create new file Upload files Find file History

k8s-merge-robot committed on GitHub Merge pull request #52628 from richardalberto/metadata-search-order Latest commit 3eb5ba8 2 hours ago

providers Merge pull request #52628 from richardalberto/metadata-search-order 2 hours ago

build Use buildozer to delete licenses() rules except under third_party/ 2 months ago

README.md Update docs/ URLs to point to proper locations. 1 month ago

cloud.go Remove links to GCE, AWS cloud providers from PersistentVolumeController 4 months ago


plugins.go add external cloudprovider to clearly denote the offloading off cloudp... 20 days ago

README.md

Deprecation Notice: This directory has entered maintenance mode and will not be accepting new providers. Cloud Providers in this directory will continue to be actively developed or maintained and supported at their current level of support as a longer-term solution evolves.



实现新的Internal Cloud Provider?






Inter


 andyzhangx commented 4 days ago Member + 😊 ✎ ✕

@jingxu97 could you approve this PR by"/approve no-issue", thx. Let me know if there is any other issue.

Add more commits by pushing to the `azuredisk-getdevfunc` branch on `andyzhangx/kubernetes`.

  **Some checks haven't completed yet** Hide all checks
2 pending and 12 successful checks

✓ 	pull-kubernetes-e2e-gce-bazel — Job succeeded.	Details
✓ 	pull-kubernetes-e2e-gce-etcd3 — Jenkins job succeeded.	Required Details
✓ 	pull-kubernetes-e2e-gce-gpu — Job succeeded.	Details
✓ 	pull-kubernetes-e2e-kops-aws — Jenkins job succeeded.	Required Details
✓ 	pull-kubernetes-kubemark-e2e-gce — Jenkins job succeeded.	Required Details

 **This pull request can be automatically merged by project collaborators**
Only those with [write access](#) to this repository can merge pull requests.

进行快速迭代

Proposal: refactor Cloud Provider out of Kubernetes Core

将cloud provider相关代码迁移到云厂商各自repo

计划 (<https://github.com/kubernetes/features/issues/88>)

1.7 alpha

1.8,1.9 beta

1.10 Stable: 将移除cloud provider相关代码

Standalone **Azure** cloud provider#50752

- Flexvolume插件
- External(Out-of-Tree) Provisioner
 - <https://github.com/kubernetes-incubator/external-storage>

flexvolume 插件需要实现的主要接口

init – initialize the driver

WaitForAttach - get device info

Attach - attach device to the node

MountDevice - mount (&format device) in file system in the node

UnmountDevice - unmount device in file system in the node

Detach - detach device from the node

如何使用

在各个k8s node及master指定目录安装flexvolume 插件

主办：



使用nfs flexvolume插件

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-nfs
spec:
  containers:
    - name: nginx-nfs
      image: nginx
      volumeMounts:
        - name: test
          mountPath: /data
  volumes:
    - name: test
      flexVolume:
        driver: "k8s/nfs"
        fsType: "nfs"
        options:
          server: "172.16.0.25"
          share: "share"
```

需要实现的主要接口

provision

delete

main() – run as a pod in k8s

在kubernetes集群中以POD方式运行

1. 基于out-of-tree provisioner code生成docker image
2. 在k8s集群中运行一个基于该provisioner image的POD
3. 为POD设置RBAC权限访问k8s集群
4. 使用out-of-tree provisioner: 生成StorageClass和PVC

主办：



- Internal(in-tree) Provisioner
- Flexvolume插件
- External(Out-of-Tree) Provisioner
 - <https://github.com/kubernetes-incubator/external-storage>

Q&A

Thank you for your time

主办：

