



# Kubernetes Storage Architecture and Evolution

任玉泉

# Content

- kubernetes storage overview
- kubernetes storage implementation
- kubernetes storage usage evolution
- kubernetes storage future features
- Q&A



架构迎接未来变化

IAS2017 • NANJING

# Kubernetes Design Principles

## Content

- **Declarative > imperative:** State your desired results, let the system actuate
- **Control loops:** Observe, rectify, repeat
- **Simple > Complex:** Try to do as little as possible
- **Modularity:** Components, interfaces, & plugins
- **Legacy compatible:** Meet users where they are, requiring apps to change is a non-starter
- **Open > Closed:** Open Source, standards, REST, JSON, etc.



# Storage Architecture Overview

## Master

Kube API Server PV PVC SC Node Pod

### Kube Controller Manager

#### Attach/Detach Controller Expand Controller

Populator

PV/PVC/  
Node/Pod  
Informers

Desired State of World

Reconciler

Actual State of World

#### Operation Executor

GoRoutineMap

Volume  
Plugins

Attacher

Mounter

Provisioner

Detacher

Unmounter

Deleter

#### PV/PVC Controller

PV/PVC/SC  
Informers

Volume  
Queue

Claim  
Queue

Volume  
Worker

Claim  
Worker

GoRoutineMap

## Node

Kubelet

Pod Manager

Volume  
Manager

Populator

Desired State of World

Reconciler

Actual State of World

#### Operation Executor

GoRoutineMap

Volume  
Plugins

Provisioner

Deleter

Mounter

Attacher

Unmounter

Detacher

Containers

Mounts

Devices

Root FS

Storage Provider Control Plane

# Kubernetes Supported Storage

## Persistent

- GCE Persistent Disk
- AWS Elastic Block Store
- Azure File Storage
- Azure Data Disk
- iSCSI
- Flocker
- NFS
- vSphere
- GlusterFS
- Ceph File and RBD
- Cinder
- Quobyte Volume
- FibreChannel
- VMWare Photon PD
- Portworx
- Dell EMC ScaleIO
- StorageOS

## Ephemeral

- Empty dir (and tmpfs)
- Expose Kubernetes API
  - Secret
  - ConfigMap
  - DownwardAPI

## New

- Local Storage



# Kubernetes Storage Implementation

- Volume Plugin Interface
- Kubelet Volume Manager
- Attach/Detach Controller
- PV/PVC Controller
- ExpandVolume Controller



架构迎接未来变化

IAS2017 • NANJING

# Kubernetes Storage Implementation

## Volume Plugin Interface

- Golang packages in core  
Kubernetes repository
  - `kubernetes/pkg/volume/`
- Implement golang interfaces
  - Mounter
  - Unmounter
  - Optionally
    - Attacher
    - Detacher
    - Provisioner
    - Deleter
    - Recycler

# Kubernetes Storage Implementation

## Volume Plugin Interface

- Mounter/Unmounter Interface
  - Make data source (volume, block device, network share, or something else) available as a directory on host's root FS.
  - Directory then mounted into pods by kubelet
  - Methods always called from node (Kubelet binary)
- Methods
  - `SetUpAt(dir, ...)`
  - `TearDownAt(dir)`
  - ...





# Kubernetes Storage Implementation

## Volume Plugin Interface

- Attacher/Detacher Interface
  - Make block device available on specified host.
  - Attach & VolumesAreAttached methods called from master (kube controller binary).
- Methods
  - `Attach(spec, nodeName)`
  - `VolumesAreAttached(specs, nodeName)`
  - `WaitForAttach(spec, devicePath, timeout)`
  - `MountDevice(spec, devicePath, deviceMountPath)`
  - `UnmountDevice(deviceMountPath)`
  - ...

# Kubernetes Storage Implementation

## Volume Plugin Interface

- Provisioner/Deleter Interface
  - Create and delete new pieces of physical storage and the k8s PV object to represent it.
  - Methods called from master (kube controller binary).
- Methods
  - `Provision()`
  - `Delete()`



# Kubernetes Storage Implementation

## Volume Plugin Interface

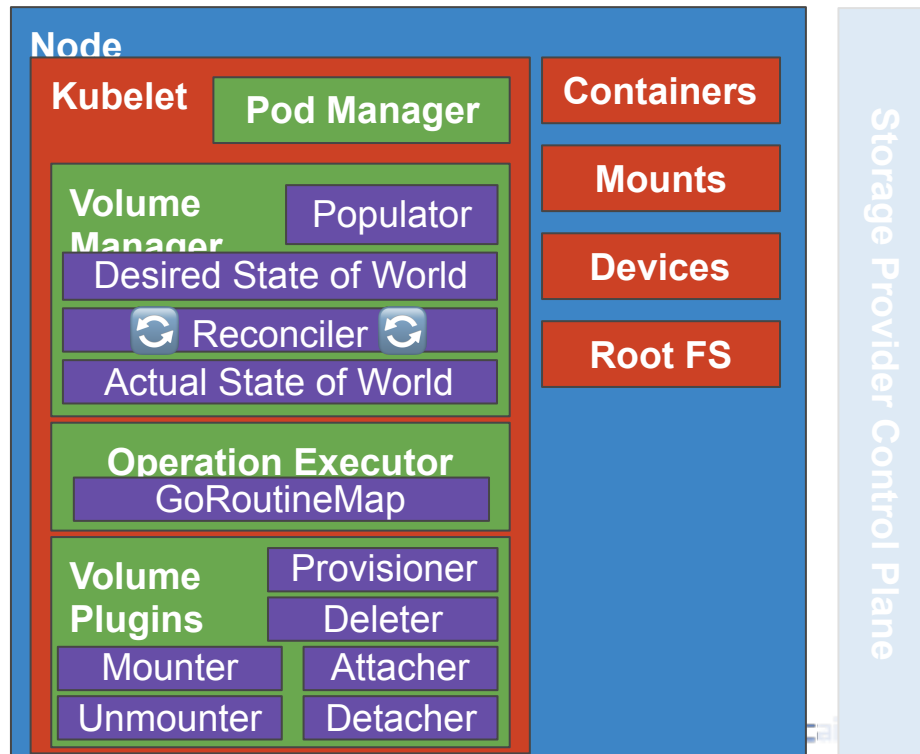
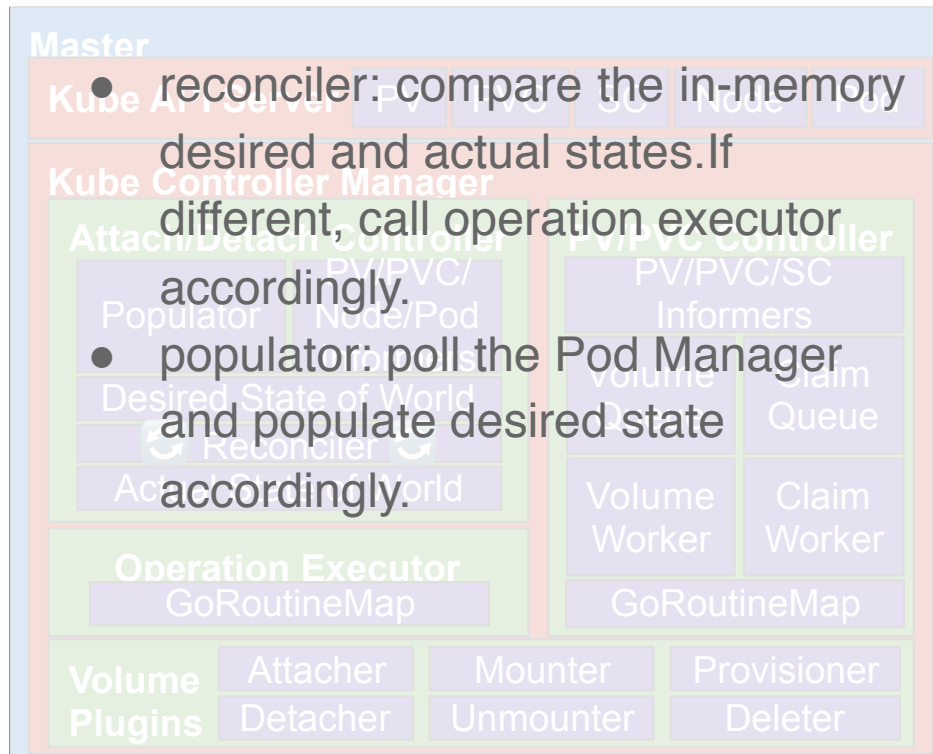
- Take cinder as an example
  - `create cinder volume (provision)`
  - `attach to instance`
  - `mount device (/var/lib/kubelet/plugins/kubernetes.io/cinder/mounts/cinder-volume-id)`
  - `mounted to pod volume dir (/var/lib/kubelet/pods/{podUID}/volumes/kubernetes.io~cinder/{outerVolumeSpecName}/)`



架构迎接未来变化  
IAS2017 • NANJING

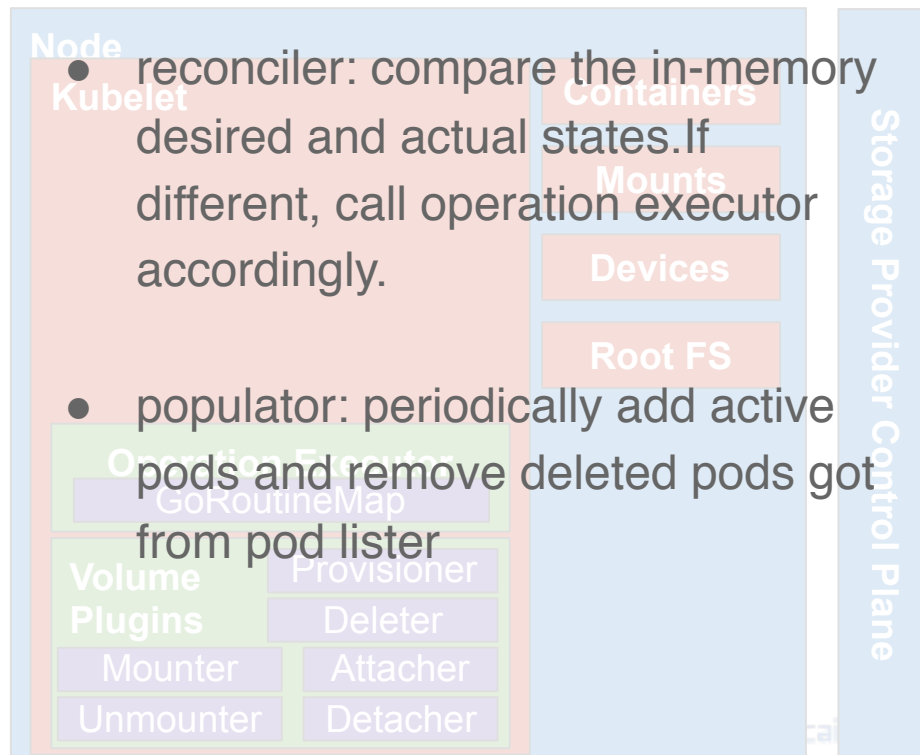
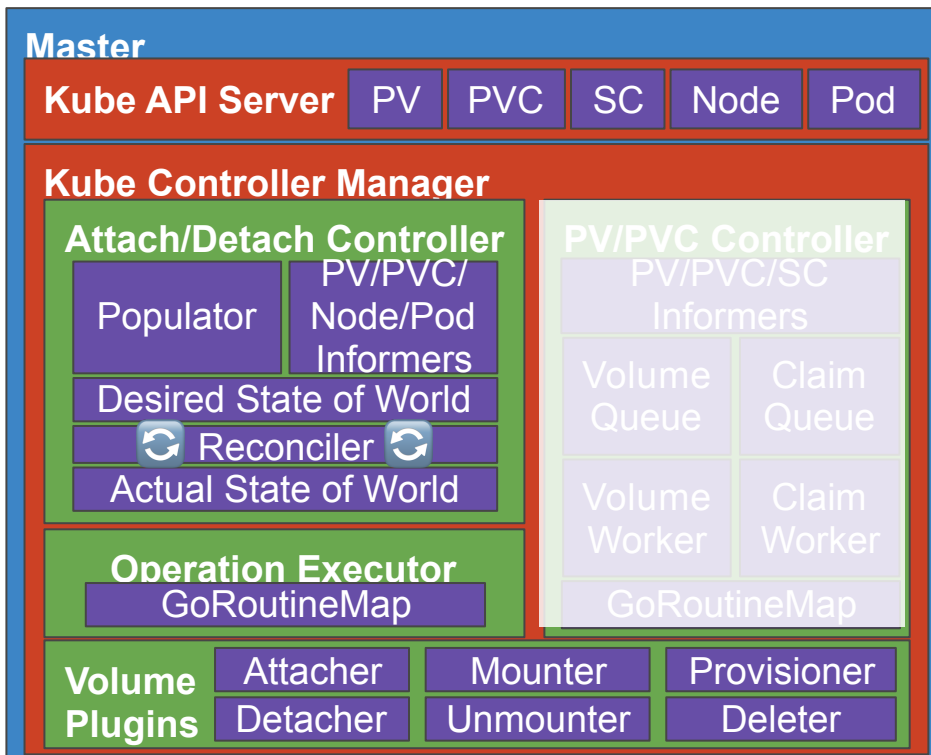
# Kubernetes Storage Implementation

## Kubelet Volume Manager



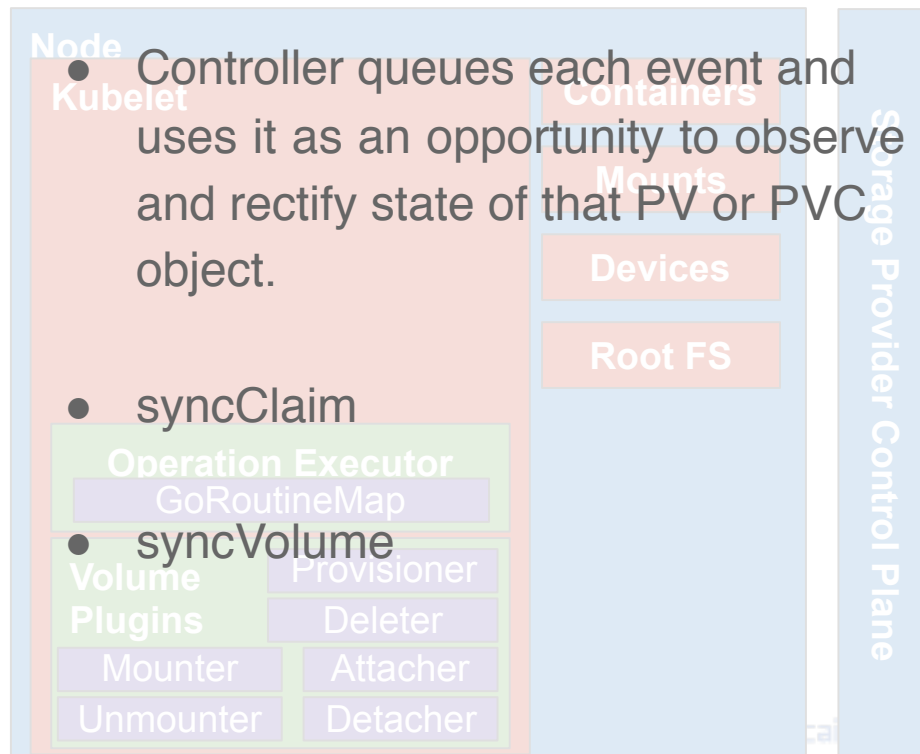
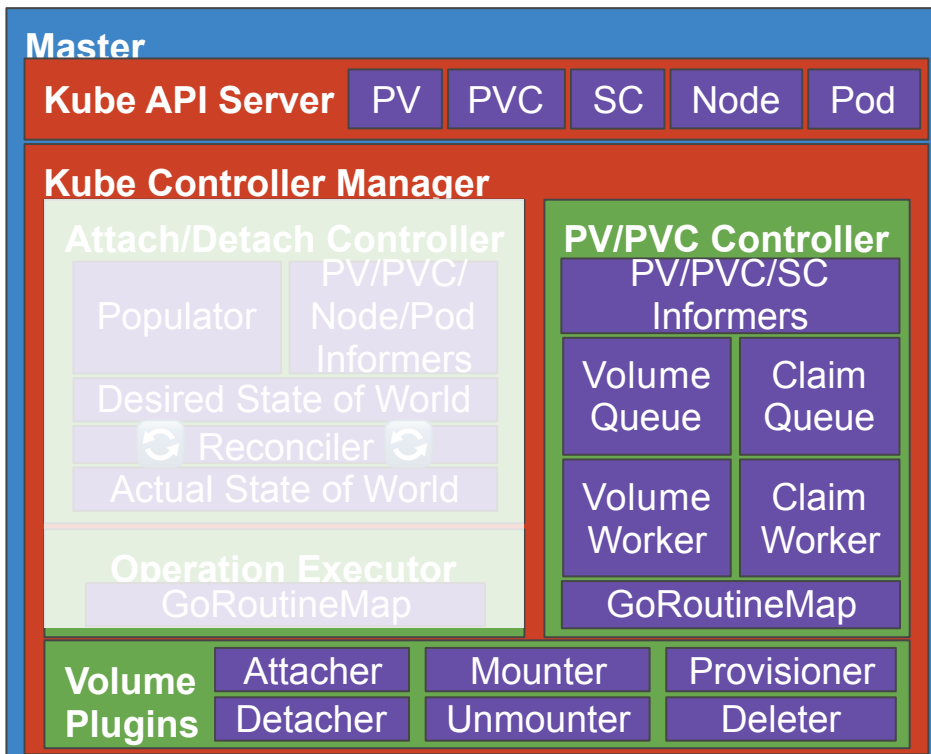
# Kubernetes Storage Implementation

## Attach/Detach Controller



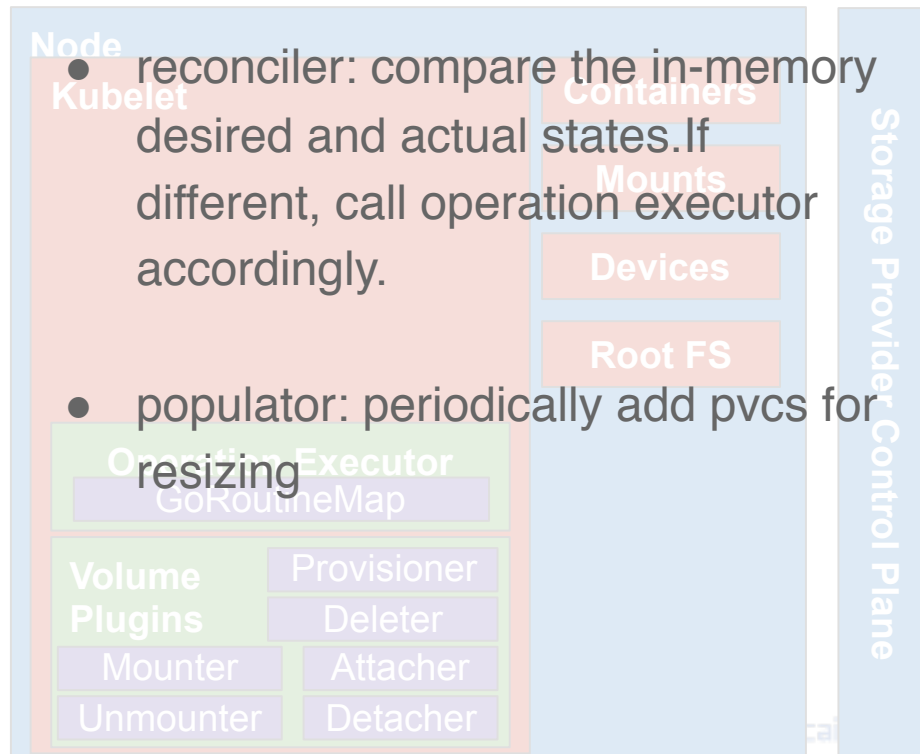
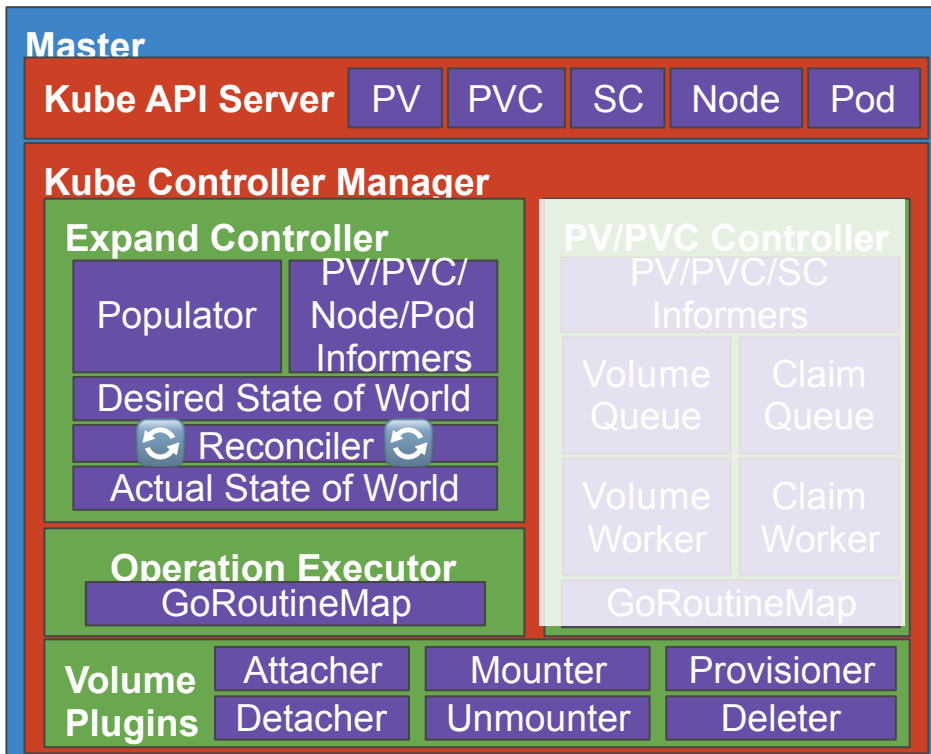
# Kubernetes Storage Implementation

## PV/PVC Controller



# Kubernetes Storage Implementation

## ExpandVolume Controller



# Kubernetes Storage Usage evolution

## Direct Access:

- Directly write volume details in Pod configuration
- Same approach for all kinds of volumes, i.e. persistent, local, ephemeral, etc

```
kind: Pod
apiVersion: v1
metadata:
  name: mypod
spec:
  containers:
    - name: nginx
      image: nginx:1.13
      volumeMounts:
        - mountPath: "/var/www/html"
          name: mypath
    - name: busybox
      image: busybox:1.26
      command: ["sh", "-c", "sleep 12000"]
      volumeMounts:
        - mountPath: "/var/www/html"
          name: mypath
  volumes:
    - name: mypath
      hostPath:
        path: /tmp/data
```

Host Path

NFS


```
apiVersion: v1
kind: Pod
metadata:
  name: pod-nfs
spec:
  containers:
    - name: nginx
      image: nginx:1.13
      volumeMounts:
        - name: storage
          mountPath: /data/storage
        - name: scratch
          mountPath: /data/scratch
  volumes:
    - name: storage
      nfs:
        path: /var/export1
        server: 192.168.44.44
    - name: scratch
      nfs:
        path: /var/export2
        server: 192.168.44.44
```



# Kubernetes Storage Usage evolution

## Direct Access:

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-nfs
spec:
  containers:
  - name: nginx
    image: nginx:1.13
    volumeMounts:
    - name: storage
      mountPath: /data/storage
    - name: scratch
      mountPath: /data/scratch
  volumes:
  - name: storage
    nfs:
      path: /var/export1
      server: 192.168.44.44
  - name: scratch
    nfs:
      path: /var/export2
      server: 192.168.44.44
```



### Observation:

- Pod is created and scheduled on a Node
  - scheduling is **independent** of volume
- Kubelet has built-in plugin libraries
  - one for each supported volume type
- Two **existing** NFS volumes are attached to Pod
  - no provisioning
  - no configuration knob
- More

# Kubernetes Storage Usage evolution

- Root problem with direct access
  - Tight coupling between setting up storage and request/use storage
- Solution
  - Add another layer which separate the complexity: admin sets up storage, user requests storage




Image Source: Google



拥抱迎接未来变化

IAS2017 • NANJING

# Kubernetes Storage Usage evolution

- Admin <- PersistentVolume (PV)
  - Persistent volume represents a schedulable, requestable storage identity
  - Can be networked storage, local storage, etc
- User <- PersistentVolumeClaim (PVC)
  - Claim volumes of specific size and modes




Image Source: Google

# Kubernetes Storage Usage evolution

```
kind: Pod
apiVersion: v1
metadata:
  name: mypod
spec:
  containers:
    - name: nginx
      image: nginx:1.13
      volumeMounts:
        - mountPath: "/var/www/html"
          name: mypd
  volumes:
    - name: mypd
      gcePersistentDisk:
        pdName: disk-1
        fsType: ext4
```

Any problems ?



```
kind: Pod
apiVersion: v1
metadata:
  name: mypod
spec:
  containers:
    - name: nginx
      image: nginx:1.13
      volumeMounts:
        - mountPath: "/var/www/html"
          name: mypd
  volumes:
    - name: mypd
      persistentVolumeClaim:
        claimName: myclaim
```

# Kubernetes Storage Usage evolution

- StorageClass is an API object created by admin to enable dynamic provisioning
  - Create PersistentVolume on request
  - Allow more configuration parameters

```
apiVersion: storage.k8s.  
kind: StorageClass  
metadata:  
  name: standard  
  labels:  
    addonmanager.kuberne  
  annotations:  
    storageclass.beta.ku  
provisioner: k8s.io/mini
```

```
kind: StorageClass  
apiVersion: storage.k8s.io/v1  
metadata:  
  name: fast  
provisioner: kubernetes.io/rbd  
reclaimPolicy: retain  
parameters:  
  monitors: 10.16.153.105:6789  
  adminId: kube  
  adminSecretName: ceph-secret  
  adminSecretNamespace: kube-system  
  pool: kube  
  userId: kube  
  userSecretName: ceph-secret-user  
  fsType: ext4  
  imageFormat: "2"  
  imageFeatures: "layering"
```

# Kubernetes Storage Usage evolution

Future is coming!






```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: standard
labels:
  addmanager.kubernetes.io/mode: Reconcile
annotations:
  storageclass.beta.kubernetes.io/is-default-class: "true"
provisioner: k8s.io/minikube-hostpath
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  storageClassName: standard
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 8Gi
```

Watch All new Claims, for each one, find its StorageClass based on spec.storageClassName, then provision new PV if class.provisioner match my name.

# Kubernetes Storage Usage evolution

- Evolution Path



# Kubernetes Storage Future Features

## Content

- Local ephemeral storage
- PVC resize
- Local persistent storage



架构迎接未来变化

IAS2017 • NANJING



# Kubernetes Storage Future Features

## Content

- Local ephemeral storage

```
apiVersion: v1
kind: Node
metadata:
  name: foo
status:
  capacity:
    ephemeral-storage: "100Gi"
  allocatable:
    ephemeral-storage: "100Gi"
```

```
apiVersion: v1
kind: pod
metadata:
  name: foo
spec:
  containers:
    - name: fooa
      image: fooa
      resources:
        requests:
          ephemeral-storage: "10Gi"
        limits:
          ephemeral-storage: "10Gi"
    - name: foob
      image: foob
      resources:
        requests:
          ephemeral-storage: "20Gi"
        limits:
          ephemeral-storage: "20Gi"
      volumeMounts:
        - name: myEmptyDir
          mountPath: /mnt/data
  volumes:
    - name: myEmptyDir
      emptyDir:
        sizeLimit: "5Gi"
```

# Kubernetes Storage Future Features

- PVC resize

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
  namespace: default
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 8Gi
  storageClassName: standard
  volumeName: pv-hostpath
status:
  accessModes:
    - ReadWriteMany
  capacity:
    storage: 10Gi
  phase: Bound
```



```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
  namespace: default
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 20Gi
  storageClassName: standard
  volumeName: pv-hostpath
status:
  accessModes:
    - ReadWriteMany
  capacity:
    storage: 10Gi
  phase: Bound
```

# Kubernetes Storage Future Features

## Content

- Local persistent storage

```
kind: PersistentVolume
apiVersion: v1
metadata:
  name: local-pv
  labels:
    kubernetes.io/hostname: node-1
  annotations:
    volume.alpha.kubernetes.io/node-affinity: >
    {
      "requiredDuringSchedulingIgnoredDuringExecution": {
        "nodeSelectorTerms": [
          {
            "matchExpressions": [
              {
                "key": "kubernetes.io/hostname",
                "operator": "In",
                "values": ["kube-node-1"]
              }
            ]
          }
        ]
      }
    }
spec:
  capacity:
    storage: 10Gi
  local:
    path: /tmp/local-pv
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Delete
  storageClassName: local-fast
```

# End

Content

# Thank you !



架构迎接未来变化

IAS2017 • NANJING