











### **Outline**

#### Overview

- Monitoring for containers
- Monitoring in a distributed system
- cAdvisor

#### **Application Metrics**

- In cAdvisor
- Plumbing through a cluster

#### Future Work

- Heapster/Kubedash
- Templates



### Monitoring in three parts

#### Collection

- collecting data and metrics
- making it discoverable
- plumbing it to higher-level systems

#### **Processing**

- Ingesting, aggregation
- Analytics

### Managing

- Actions based on signals
- Alerts. Pagers!



### Monitoring in three parts

#### Collection

- collecting data and metrics
- making it discoverable
- plumbing it to higher-level systems

#### **Processing**

- Ingesting, aggregation
- Analytics

### Managing

- Actions based on signals
- Alerts. Pagers!



#### Collection in servers/VMs

- Node agent
- Knows the binary to monitor
- Monitoring logic plugged into agent

#### **Moving to containers**

- Same node agent understands multiple applications.
- Applications can push data to agent.
- Sidecars
- Applications publish data. Pulled off by offhost agents.

### consumers

#### **Near-instant data**

- schedulers in cluster management tools
- load balancers
- Alerting systems

#### Slower feedback loops

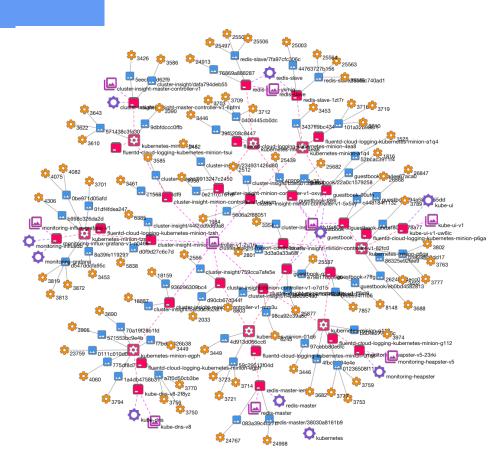
- Autoscaling / Autoupdaters
- <u>CPI</u><sup>2</sup>

#### **Historical Data**

Offline analysis.

# Monitoring in distributed environment

- Hosts are invisible
- Containers can hop around
- Monitoring needs to track and correlate multiple containers



### In Action!



### **c**Advisor

Analyzes resource usage and performance characteristics of running containers

Native Docker support, work with LXC and any other container format



**Knows containers** deeply and monitors their performance

github.com/google/cadvisor

# **Running cAdvisor**

```
docker run
  --volume=/:/rootfs:ro
  --volume=/var/run:/var/run:rw
  --volume=/sys:/sys:ro
  --volume=/var/lib/docker/:/var/lib/docker:ro
  --publish=8080:8080
  --detach=true
  --name=cadvisor
 google/cadvisor:latest
```

### **API**

Node and container spec

http://host:8080/api/v2.0/machine

http://host:8080/api/v2.0/spec/redis?type=docker

**Hierarchical Container stats:** 

http://host:8080/api/v2.0/stats/nginx?type=docker

http://host:8080/api/v2.0/stats?recursive=true

Others:

/summary, /events, /storage, /attributes, /ps

# Storage Backends

















Coming soon ...







# **Application Metrics**

Every container is packaged with its monitoring data.

Monitoring moves with the container.

Use container composability to stack up metrics from all layers.

Use metadata in image or runtime (Docker labels!) to configure monitoring

# **Configuring containers**

```
FROM redis
ADD redis config.json /var/cadvisor/redis config.json
LABEL io.cadvisor.metric.redis="/var/cadvisor/redis config.json"
In cAdvisor
read Labels "io.cadvisor.metric.*"
read /rootfs/proc/<pid>/root/<config path>
```

# **Monitoring Configuration**

#### Holds metadata about metrics

- Endpoint (Location to collect metrics from)
- Name of metric
- Type (Counter, Gauge, ...)
- Data Type (int, float)
- Units (kbps, seconds, count)
- Polling Frequency
- Regexps (Regular expressions to be used to collect a metric)

# **Sample Configurations**

```
Get all prometheus metrics:
      "endpoint": "http://localhost:9100/metrics",
Get selected prometheus metrics:
      "endpoint": "http://localhost:8000/metrics",
      "metrics_config" : [
            { "scheduler_binding_latency",
              "scheduler_e2e_scheduling_latency",
              "scheduling_algorithm_latency"
```

```
{
      "endpoint": "http://localhost:8000/nginx_status",
      "metrics_config" : [
            { "name" : "activeConnections",
              "metric_type" : "gauge",
              "units": "number of active connections",
              "data_type" : "int",
              "polling_frequency": 10,
              "regex": "Active connections: ([0-9]+)"
            { "name" : "reading",
              "metric_type" : "gauge",
              "units": "number of reading connections",
              "data_type" : "int",
              "polling_frequency": 10,
              "regex": "Reading: ([0-9]+).*"
```

# **App Metrics in action**







### **API**

Endpoint for custom metrics:

http://localhost:8080/api/v2.0/appmetrics/containerName

Application metrics being collected can be discovered from the spec:

http://localhost:8080/api/v2.0/spec/containerName

Regular stats api also reports application metrics:

http://localhost:8080/api/v2.0/stats/containerName

### Kubernetes

Open-source Container orchestration from Google

Declarative pattern for managing containers Physical hosts abstracted out as resources

Inspired and informed by Borg

Kubernetes UI powered by cAdvisor

github.com/kubernetes/kubernetes



# **App Metrics for Kubernetes**

All kubernetes components run in containers

All system services on node runs in containers

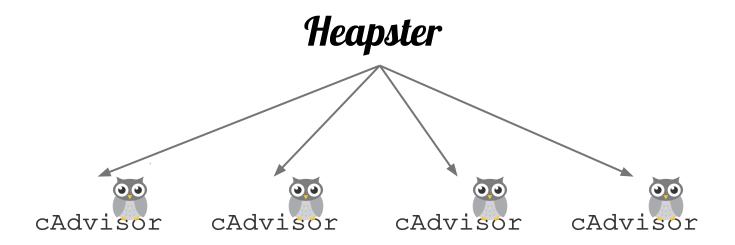
All kubernetes nodes run cAdvisor (built into kubelet)

All kubernetes components expose prometheus metrics

THE PARTY OF THE P

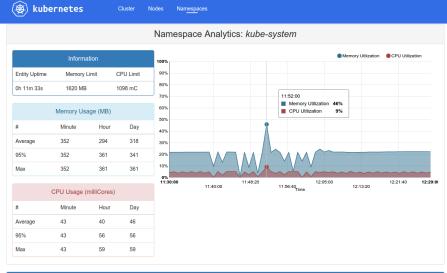
App-metrics for kubernetes components!

# Heapster



github.com/kubernetes/heapster

### Kubedash



Pods in kube-system		
Name \$	Memory Usage (MB) ▼	CPU Usage (mC) ÷
fluentd-cloud-logging-kubernetes-minion-nxh2	89	9
fluentd-cloud-logging-kubernetes-minion-6rpk	67	12
fluentd-cloud-logging-kubernetes-minion-k43o	62	8
fluentd-cloud-logging-kubernetes-minion-lsuf	61	8
monitoring-heapster-v6-ug03w	41	0
monitoring-influx-grafana-v1-mog7z	19	0
kube-dns-v8-74o3k	11	6
kube-ui-v1-utyyr	1	0

github.com/kubernetes/kubedash

# **Templates**

Add templates for applications that have stable stats API

```
LABEL io.cadvisor.metric.type=redis
```

Infer monitoring information

- Lookup ports through docker inspection
- hit known endpoints (e.g.: /metrics for prometheus)
- overrides through config

# **Tags**

Adding tags to specific metrics Convey metric intent to processors eg. Autoscalers

# **Ongoing work**

Endpoints
Storage drivers
Standard config syntax
Automagic

Plumbing through heapster/kubedash/kubernetes/...

# Thank you!

Rohit Jnagal jnagal@google Anushree Narasimha anushree.bnp@gmail

cAdvisor <u>github.com/google/cadvisor</u>

kubernetes kubernetes.io

irc #google-containers