


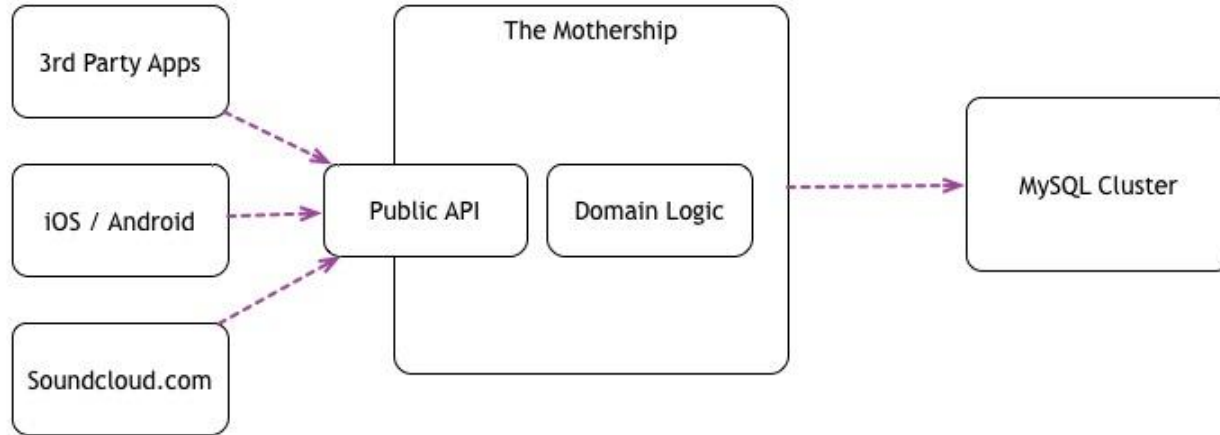
# Monitoring a **Kubernetes-backed** microservice architecture with **Prometheus**

Björn “Beorn” Rabenstein — *SoundCloud*

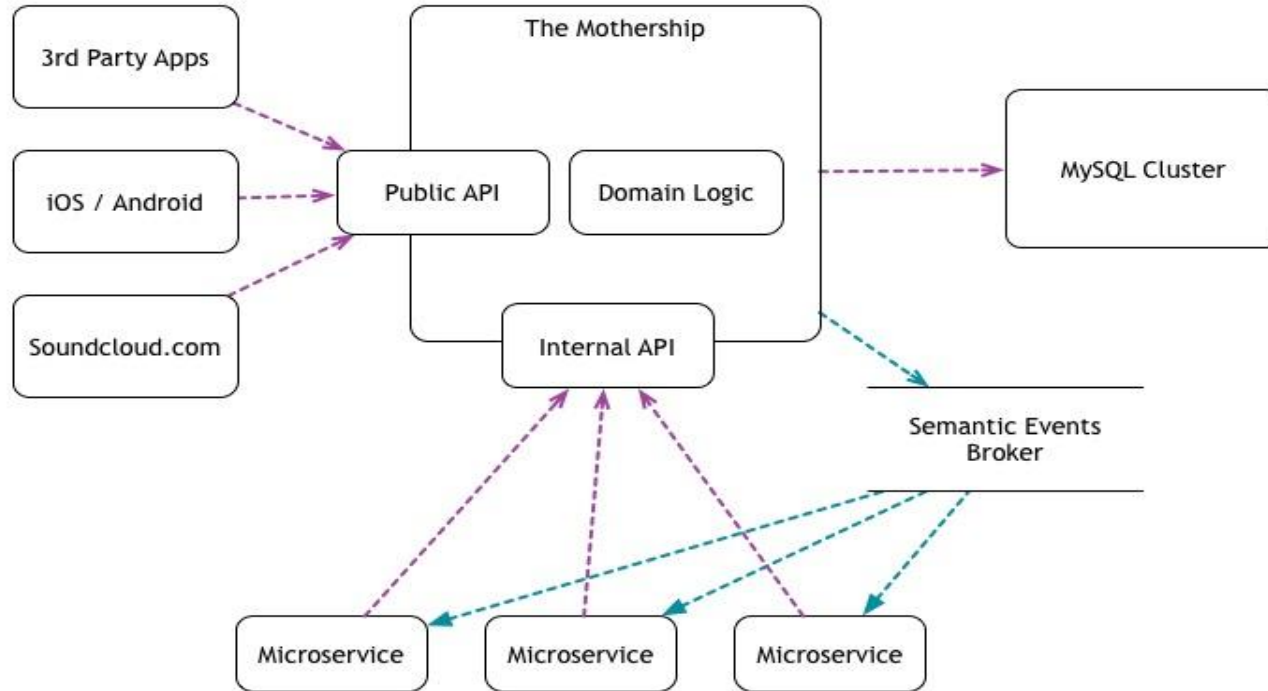
Fabian Reinartz — *CoreOS*



# SoundCloud 2012 – from monolith ...



... to microservices



# Orchestration needed

Run containers in a cluster...

In-house innovation: *Bazooka* – PaaS, Heroko style.

Problems:

- Only 12-factor apps (stateless etc.).
- Limited resource isolation.
- No sidecars.
- Maturity.

Meanwhile, the open-source world has evolved...



## THE TWELVE-FACTOR APP






KUBERNETES

# Kubernetes

- inspired by Google's Borg
- not Borg



## Today:



14:45 - 15:35

VICTORIA SUITE



**Container Orchestration with Kubernetes**

[Peter Rossbach](#), *bee42 solutions GmbH*

## Tomorrow:



09:00 - 09:50

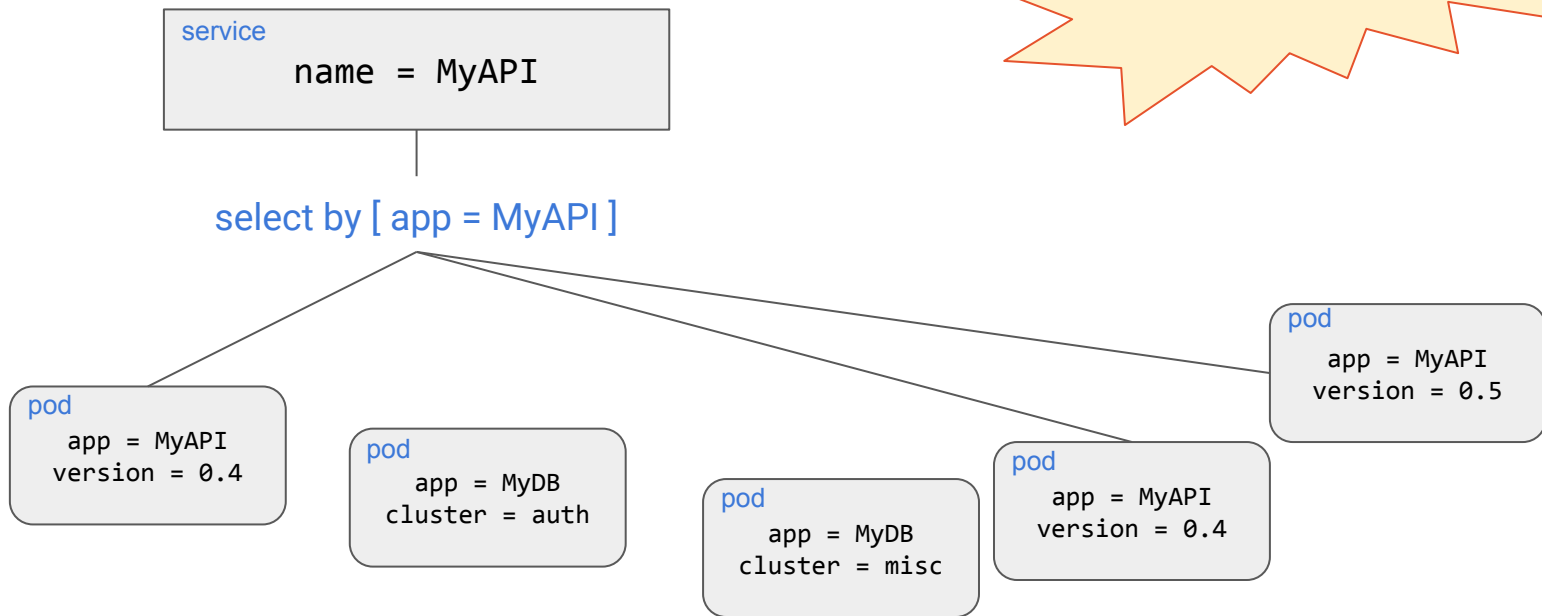
ALBERT SUITE



**Java-based microservices, containers, Kubernetes – how to**


[Ray Tsang](#), *Google*

# Labels



# Monitoring at SC 2012

**Nagios**<sup>®</sup>





# Monitoring challenges

- A lot of traffic to monitor
  - Monitoring traffic should not be proportional to user traffic
- Way more targets to monitor
  - One host can run many containers
- And they constantly change
  - Deploys, scaling, rescheduling unhealthy instances ...
- Need a fleet-wide view.
  - What's my overall 99th percentile latency?
- Still need to be able to drill down for troubleshooting
  - Which instance/endpoint/version/... causes those errors I'm seeing?
- Meaningful alerting
  - Symptom-based alerting for pages, cause-based alerting for warnings
  - See Rob Ewaschuk's *"My philosophy on alerting"* <https://goo.gl/2vrpSO>

# Monitor everything, all levels, with the same system

Level	What to monitor (examples)	What exposes metrics (example)
Network	Routers, switches	SNMP exporter
Host (OS, hardware)	Hardware failure, provisioning, host resources	Node exporter
Container	Resource usage, performance characteristics	cAdvisor
Application	Latency, errors, QPS, internal state	Your own code
Orchestration	Cluster resources, scheduling	Kubernetes components

*“Obviously the solution to all our problems with everything forever, right?”*

Benjamin Staffin, Fitbit Site Operations



P R O M E T H E U S

# Prometheus

- inspired by Google's Borgmon
- not Borgmon
- initially developed at SoundCloud, open-source from the beginning
- public announcement early 2015
- collects metrics at scale via HTTP (think: yet another client of your microservice)
- thousands of targets, millions of time series, 800k samples/s, no dependencies
- easy to scale

# Features – multi-dimensional data

***Labels are the  
new hierarchies!***

`http_requests_total{instance="web-1", path="/index", status="200", method="GET"}`

`http_requests_total{instance="web-1", path="/index", status="404", method="POST"}`

`http_requests_total{instance="web-3", path="/index", status="200", method="GET"}`

`#metrics x #values(instance) x #values(path) x #values(status) x #values(method)`

► millions of time series

# Features – powerful query language

The questions to ask are often not known beforehand.

*The 3 path-method combinations with the highest number of failing requests?*

```
topk(3, sum by(path, method) (  
    rate(http_requests_total{status=~"5.."}[5m])  
))
```

*The 99th percentile request latency by request path?*

```
histogram_quantile(0.99, sum by(le, path) (  
    rate(http_requests_duration_seconds_bucket[5m])  
))
```

# Features – powerful query language

```
topk(3, sum by(path, method) (  
    rate(http_requests_total{status=~"5.."}[5m])  
))
```

{path="/api/comments", method="POST"}	105.4
{path="/api/user/:id", method="GET"}	34.122
{path="/api/comment/:id/edit", method="POST"}	29.31

# Features – easy instrumentation

```
from prometheus_client import start_http_server, Histogram

# Create a metric to track time spent and requests made.
REQUEST_TIME = Histogram('request_processing_seconds', 'Time spent processing request')


# Decorate function with metric.
@REQUEST_TIME.time()
def process_request(t):
    # do work ...
    return

start_http_server(8000)
```




# Integrations (selection)





KUBERNETES  
PROMETHEUS

O'REILLY®




# Site Reliability Engineering

HOW GOOGLE RUNS PRODUCTION SYSTEMS


Edited by Betsy Beyer, Chris Jones,  
Jennifer Petoff & Niall Murphy

# Three questions

How to monitor  
services running  
on Kubernetes  
with Prometheus?



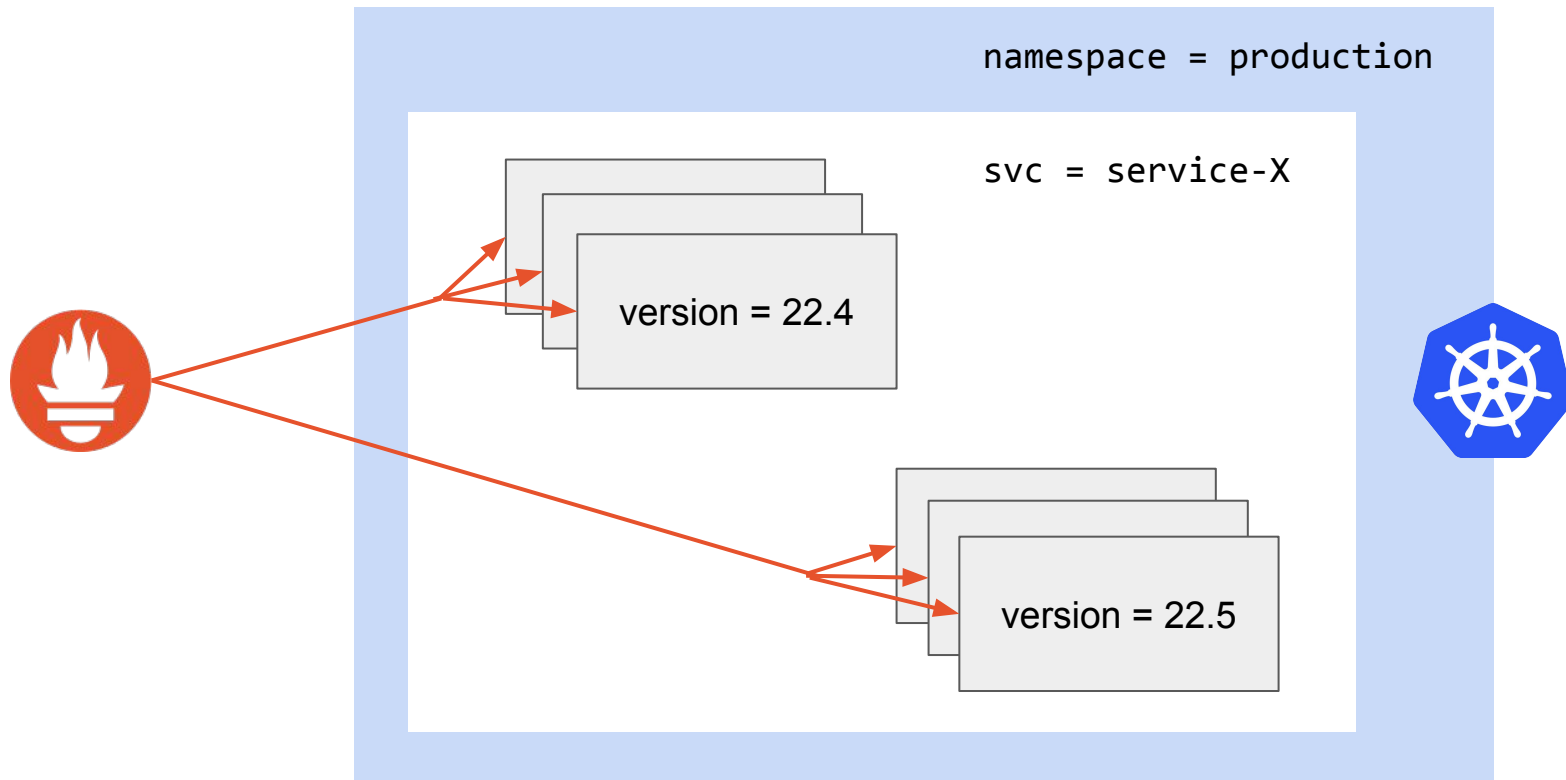
How to monitor  
Kubernetes and  
containers with  
Prometheus?



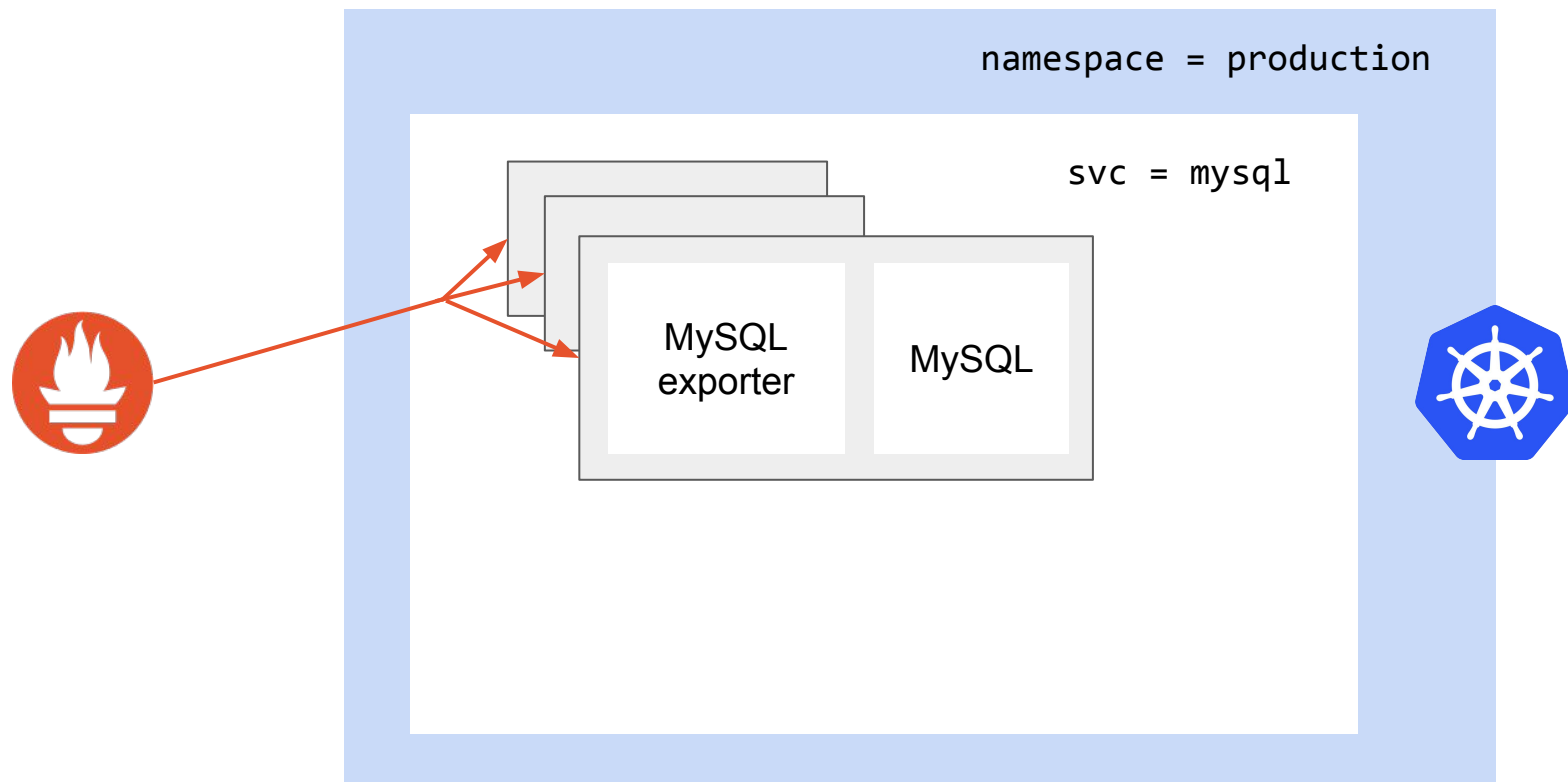
How to run  
Prometheus on  
Kubernetes?




# Monitoring Services



# Monitoring Services via Exporters



# Monitoring Kubernetes



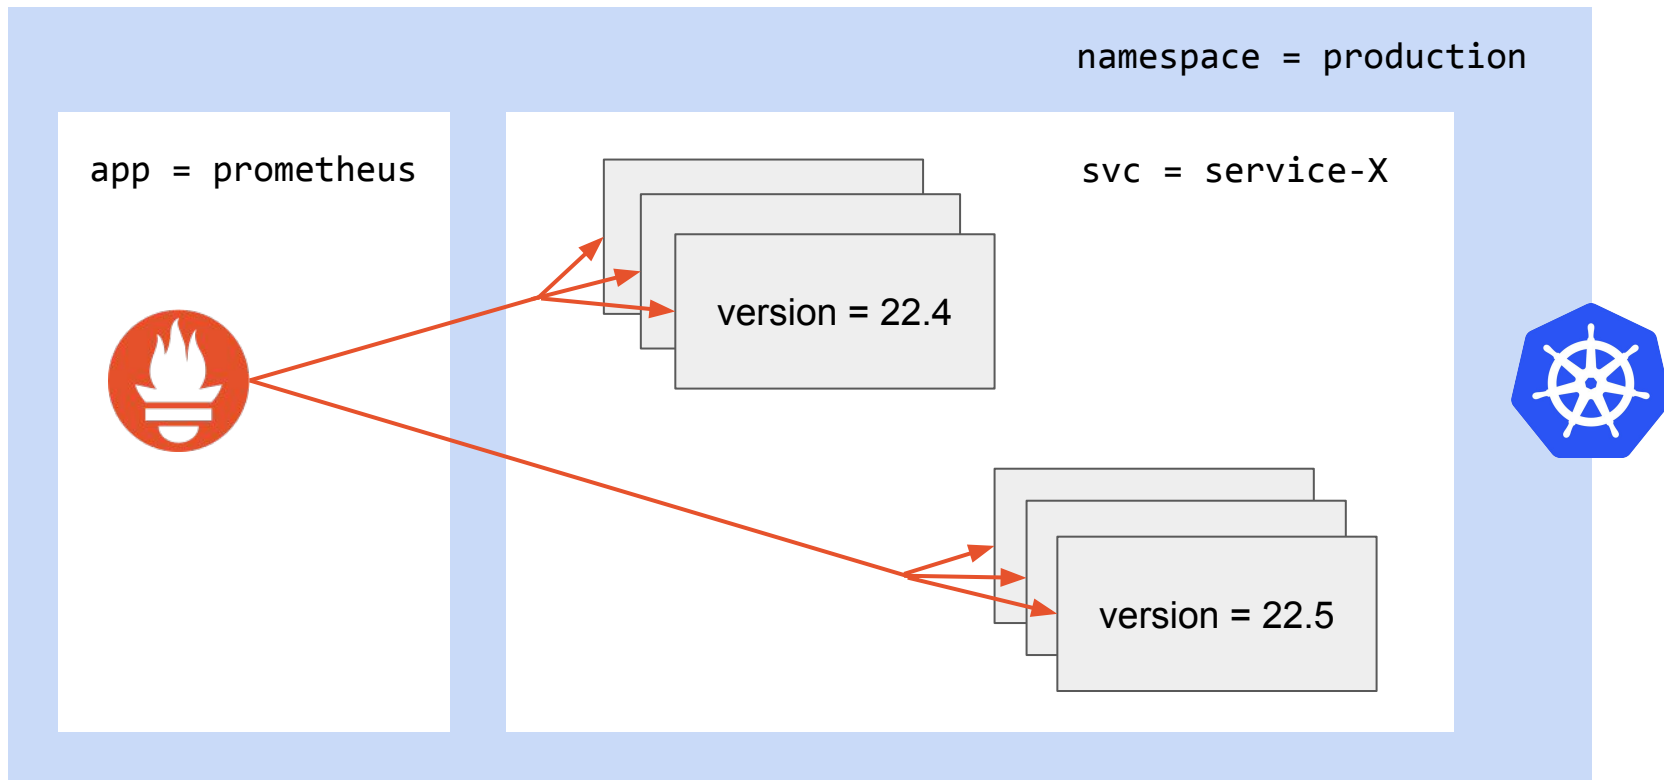
# Running Prometheus on Kubernetes

- So far: Prometheus ran outside of cluster
  - Pod IPs must be routable
  - Conventional deployment (Chef, Puppet, ...)
  - Service discovery needs authentication
- To run Prometheus inside of cluster:

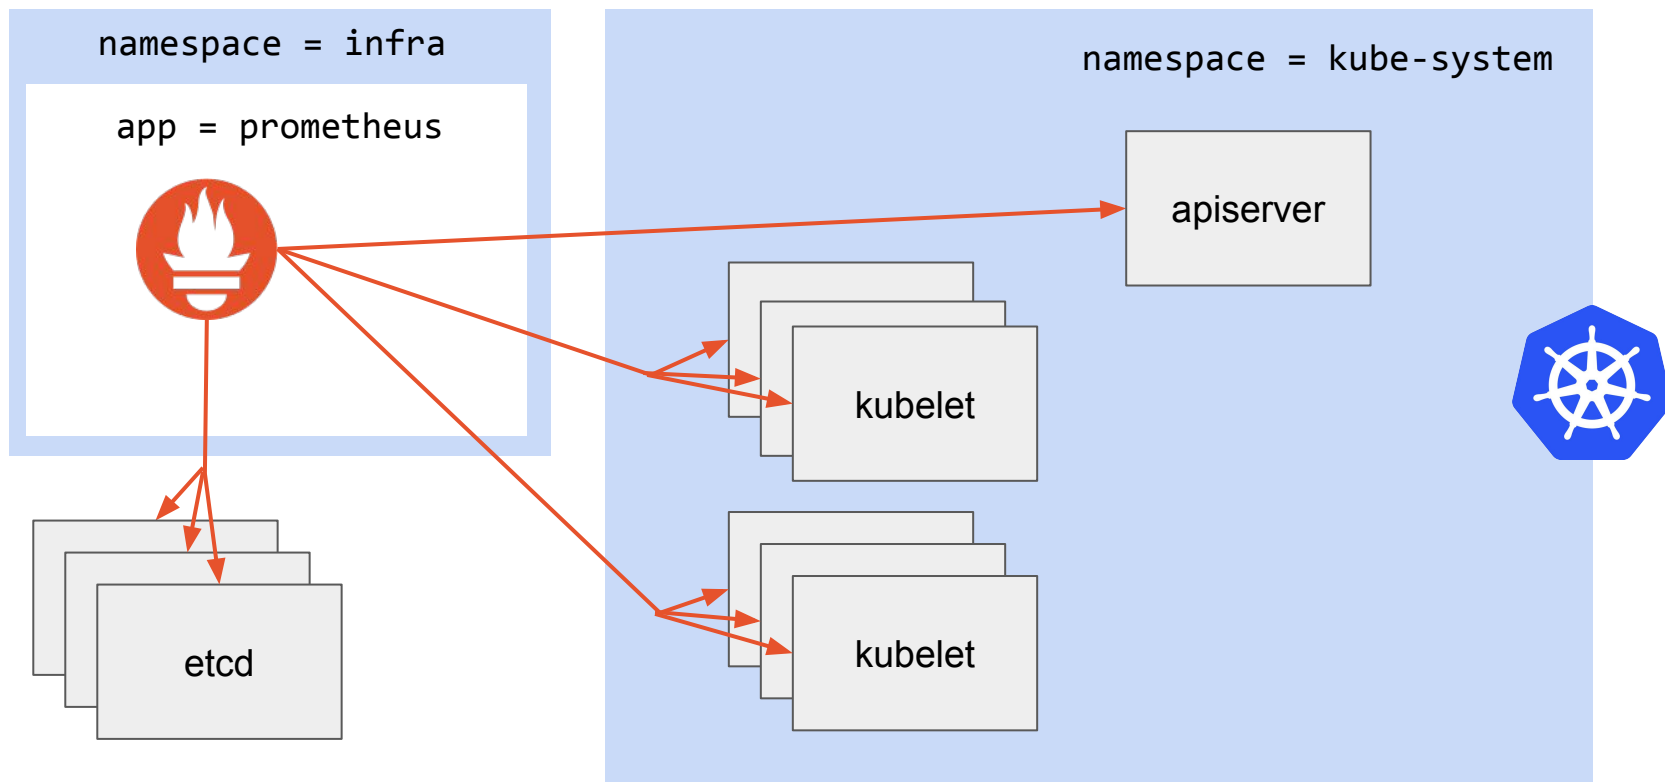
```
kubectl run --image="quay.io/prometheus/prometheus:0.18.0" prometheus
```



# Monitoring Services



# Monitoring Kubernetes



# What about storage?

A) None

B) Network/Cloud volumes

C) Host volumes



# The end

